| | |
|---|---|
| **Status** | Finished |
| **Started** | Wednesday, 13 August 2025, 7:35 PM |
| **Completed** | Wednesday, 13 August 2025, 9:01 PM |
| **Duration** | 1 hour 25 mins |
| **Marks** | 408.70/467.00 |
| **Grade** | **8.75** out of 10.00 (**87.52**%) |

Question **1**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

The Azure SQL Edge service is optimized for IoT solutions.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

The [Azure SQL Edge] service is optimized for IoT solutions.

Question **2**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Relational data uses keys to enforce relationships between different tables.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

Relational data uses [keys] to enforce relationships between different tables.

Question **3**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

The workflow for creating an analytics solution in Microsoft Power BI starts with Power BI Desktop.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

The workflow for creating an analytics solution in Microsoft Power BI starts with [Power BI Desktop.]

Question **4**

Correct

Mark 1.00 out of 1.00

You need to store data by using Azure Table storage.
What should you create first?

- a. a storage account
- b. a table
- c. an Azure Cosmos DB instance
- d. a blob container

Your answer is correct.

The correct answer is: a storage account

Question **5**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

UPDATE is an example of a Data Manipulation Language (DML) statement.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

[UPDATE] is an example of a Data Manipulation Language (DML) statement.

Question **6**

Correct

Mark 1.00 out of 1.00

Which type of non-relational data store supports a flexible schema, stores data as JSON files, and stores the all the data for an entity in the same document?

- a. document
- b. time series
- c. columnar
- d. graph

Your answer is correct.

The correct answer is:
document

Question **7**

Correct

Mark 1.00 out of 1.00

Which service is optimized for IoT scenarios that involve streaming time series data?

- a. Azure SQL Edge
- b. Azure SQL Database
- c. SQL Server on Azure Virtual Machines
- d. Azure SQL Managed Instance

Your answer is correct.

The correct answer is: Azure SQL Edge

Question **8**

Correct

Mark 2.00 out of 2.00

What are two benefits of platform as a service (PaaS) relational database offerings in Azure, such as Azure SQL Database? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [x] a. access to the latest features
- [ ] b. in-database machine learning services
- [x] c. reduced administrative effort for managing the server infrastructure
- [ ] d. complete control over backup and restore processes

Your answer is correct.

The correct answers are:
access to the latest features, reduced administrative effort for managing the server infrastructure

Question **9**

Partially correct

Mark 2.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

| | |
|---|---|
| Normalization involves eliminating relationships between database tables. | YES |
| Normalizing a database reduces data redundancy. | YES |
| Normalization improves data integrity. | YES |

Your answer is partially correct.

| **Statement** | **Yes/No** | **Reason** |
|---|---|---|
| Normalization involves eliminating relationships between database tables. | No | Normalization **organizes** data into related tables but does **not eliminate relationships**; it maintains them via foreign keys and structure. |
| Normalizing a database reduces data redundancy. | Yes | It reduces duplication by storing data efficiently in related tables, so the same data isn't repeated unnecessarily. |
| Normalization improves data integrity. | Yes | It enforces consistency and accuracy of data through structured rules and relationships, thus enhancing data integrity. |

You have correctly selected 2.

The correct answer is: Normalization involves eliminating relationships between database tables. → NO, Normalizing a database reduces data redundancy. → YES, Normalization improves data integrity. → YES

Question **10**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Database administrators apply data cleansing routines and turn data into useful information. No

Data engineers manage databases, store backup copies of data, and restore data in the event of a failure. No

Data analysts create data visuals and enable companies to make data-driven decisions. Yes

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Database administrators apply data cleansing routines and turn data into useful information. [No]
Data engineers manage databases, store backup copies of data, and restore data in the event of a failure. [No]
Data analysts create data visuals and enable companies to make data-driven decisions. [Yes]

Question **11**

Correct

Mark 3.00 out of 3.00

What are three characteristics of an Online Transaction Processing (OLTP) workload? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [ ] a. light writes and heavy reads
- [x] b. schema on write
- [ ] c. schema on read
- [x] d. normalized data
- [ ] e. denormalized data
- [x] f. heavy writes and moderate reads

Your answer is correct.

The correct answers are:
heavy writes and moderate reads,

schema on write,

normalized data

Question **12**

Partially correct

Mark 2.00 out of 3.00

- For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Azure Data Studio can be used to query an Azure SQL database from a device that runs macOS. | YES |
| Microsoft SQL Server Management Studio (SSMS) enables users to create and use SQL notebooks. | NO |
| Azure Data Studio can be used to restore a database. | NO |

Your answer is partially correct.

| Statement | Yes/No | Reason |
|---|---|---|
| **Azure Data Studio can be used to query an Azure SQL database from a device that runs macOS.** | **Yes** | Azure Data Studio is a **cross-platform tool** and can run on **macOS**, Windows, and Linux. It supports querying Azure SQL databases. |
| **Microsoft SQL Server Management Studio (SSMS) enables users to create and use SQL notebooks.** | **No** | **SSMS does not support SQL notebooks**. SQL notebooks are supported in **Azure Data Studio**. |
| **Azure Data Studio can be used to restore a database.** | **Yes** | Azure Data Studio supports **basic backup and restore operations** using T-SQL or integrated options (though SSMS offers more advanced support). |

You have correctly selected 2.

The correct answer is: Azure Data Studio can be used to query an Azure SQL database from a device that runs macOS. → YES, Microsoft SQL Server Management Studio (SSMS) enables users to create and use SQL notebooks. → NO, Azure Data Studio can be used to restore a database. → YES

Question **13**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Normalizing a database increases the throughput of writing transactions. | NO |
| Analytics systems are more normalized than transactional systems. | NO |
| Normalizing a database results in queries that require more joins. | YES |

Your answer is correct.

The correct answer is:
Normalizing a database increases the throughput of writing transactions. → NO,

Analytics systems are more normalized than transactional systems. → NO,

Normalizing a database results in queries that require more joins. → YES

Question **14**

Correct

Mark 1.00 out of 1.00

You need to perform hybrid transactional and analytical processing (HTAP) queries against Azure Cosmos DB data sources by using Azure Synapse Analytics.
What should you use?

- a. Synapse Link ✓
- b. a Synapse SQL pool
- c. Synapse pipelines
- d. Synapse Studio

Your answer is correct.

The correct answer is: Synapse Link

Question **15**

Correct

Mark 1.00 out of 1.00

You have data saved in the following format.

```
{
      "firstname": "John",
      "lastName": "Doe",
      "age": "23",
      "hobbies": [
            { "type": "Sports", "value": "Golf" },
            { "type": "Leisure", "value": "Reading" },
            { "type": "Leisure", "value": "Guitar" }
      ]
}
```

Which format was used?

- a. JSON ✓
- b. XML
- c. YAML
- d. HTML

Your answer is correct.

The correct answer is: JSON

Question **16**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Azure Synapse Analytics scales storage and compute independently | YES |
| Azure Synapse Analytics can be paused to reduce compute costs | YES |
| An Azure Synapse Analytics data warehouse has a fixed storage capacity | NO |

Your answer is correct.

The correct answer is:
Azure Synapse Analytics scales storage and compute independently → YES,

Azure Synapse Analytics can be paused to reduce compute costs → YES,

An Azure Synapse Analytics data warehouse has a fixed storage capacity → NO

Question **17**

Correct

Mark 3.00 out of 3.00

```
"customer" : {
    "first name" : "Ben",
    "last name" : "Smith",
    "address" : {
         "line 1" : "161 Azure Ln",
         "line 2" : "Palo Alto",
         "ZIP code" : "54762"
    },
    "social media": [
        {
           "service" : "twitter",
           "handle" : "@bensmith"
        },
        {
           "service" : "linkedin",
           "handle" : "bensmith"
        }
    ],
    "phone numbers": [
        {
            "type" : "mobile",
            "number" : "555-555-555"
        }
    ]
}
```

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the JSON document.

Customer is [answer choice] a root object .

Address is [answer choice] a nested object .

Social media is [answer choice] a nested array .

Your answer is correct.

The correct answer is:


```
"customer" : {
    "first name" : "Ben",
    "last name" : "Smith",
    "address" : {
         "line 1" : "161 Azure Ln",
         "line 2" : "Palo Alto",
         "ZIP code" : "54762"
    },
    "social media": [
        {
           "service" : "twitter",
           "handle" : "@bensmith"
        },
        {
           "service" : "linkedin",
           "handle" : "bensmith"
        }
    ],
    "phone numbers": [
        {
            "type" : "mobile",
            "number" : "555-555-555"
        }
    ]
}
```


Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the JSON document.

Customer is [answer choice] [a root object].

Address is [answer choice][a nested object].

Social media is [answer choice][a nested array].

Question **18**

Incorrect

Mark 0.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

The massively parallel processing (MPP) engine of Azure [redirects client connections across control nodes.] Synapse Analytics

Your answer is incorrect.

## Correct Answer:

**distributes processing across compute nodes.**

## Reason:

In Azure Synapse Analytics, the **MPP (Massively Parallel Processing)** engine breaks down large queries into smaller tasks and **distributes them across multiple compute nodes** for faster execution. This enables high performance for large-scale data processing.

- **Control node** coordinates the process but does **not** execute query steps.
- **Compute nodes** perform the actual query processing in parallel.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

The massively parallel processing (MPP) engine of Azure [distributes processing across compute nodes.] Synapse Analytics

Question **19**

Correct

Mark 1.00 out of 1.00

What is a characteristic of batch processing?

a. The data ingested during batch processing must be processed as soon as the data is received.

b. There is a noticeable time delay between ingesting data and obtaining the data processing results.

c. Large datasets must be split into batches of less than 1 GB before the data can be processed.

d. Batch processing can only process data that is structured.

Your answer is correct.

The correct answer is:
There is a noticeable time delay between ingesting data and obtaining the data processing results.

Question **20**

Correct

Mark 1.00 out of 1.00

Which scenario is an example of a streaming workload?

- a. sending telemetry data from edge devices
- b. sending transactions daily from point of sale (POS) devices
- c. sending transactions that are older than a month to an archive
- d. sending cloud infrastructure metadata every 30 minutes

Your answer is correct.

The correct answer is: sending telemetry data from edge devices

Question **21**

Correct

Mark 3.00 out of 3.00

Which three services support Apache Spark? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- a. Azure Data Explorer
- b. Azure Synapse Analytics
- c. Azure Stream Analytics
- d. Azure HDInsight
- e. Azure Databricks

Your answer is correct.

The correct answers are: Azure Synapse Analytics, Azure Databricks, Azure HDInsight

Question **22**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point

Processing salary payments once a month is an example of a batch workload. Yes

A wind turbine that sends 50 sensor readings per second is an example of a streaming workload. Yes

A home electricity meter that sends readings once a day to an energy provider is an example of a streaming workload. No

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point

Processing salary payments once a month is an example of a batch workload. [Yes]
A wind turbine that sends 50 sensor readings per second is an example of a streaming workload. [Yes]
A home electricity meter that sends readings once a day to an energy provider is an example of a streaming workload. [No]

Question **23**

Correct

Mark 1.00 out of 1.00

You need to modify a view in a relational database by adding a new column.
Which statement should you use?

a. UPDATE

b. MERGE

c. INSERT

d. ALTER

Your answer is correct.

The correct answer is:
ALTER

Question **24**

Correct

Mark 1.00 out of 1.00

What is a benefit of the Azure Cosmos DB Table API as compared to Azure Table storage?

- a. provides a higher storage capacity
- b. supports a multi-master model
- c. supports partitioning
- d. provides resiliency if an Azure region fails

Your answer is correct.

The correct answer is:
supports a multi-master model

Question **25**

Correct

Mark 1.00 out of 1.00

Which statement is an example of Data Manipulation Language (DML)?

- a. CREATE
- b. DROP
- c. ALTER
- d. INSERT

Your answer is correct.

The correct answer is: INSERT

Question **26**

Partially correct

Mark 2.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Platform as a service (PaaS) database offerings in Azure require less setup and configuration effort than infrastructure as a service (IaaS) database offerings. | YES |
| Platform as a service (PaaS) database offerings in Azure provide end users with the ability to control and update the operating system version. | NO |
| All relation and non-relational platform as a service (PaaS) database offerings in Azure can be paused to reduce costs. | YES |

Your answer is partially correct.

| Statement | Yes/No | Reason |
|---|---|---|
| Platform as a service (PaaS) database offerings in Azure require less setup and configuration effort than infrastructure as a service (IaaS). | Yes | PaaS abstracts infrastructure management, automates updates, backups, and scaling, reducing manual setup compared to IaaS. |
| Platform as a service (PaaS) database offerings in Azure provide end users with the ability to control and update the operating system version. | No | In PaaS, the **underlying OS is managed by Azure**, and users cannot directly control or update the OS version. |
| All relational and non-relational platform as a service (PaaS) database offerings in Azure can be paused to reduce costs. | No | **Not all** PaaS database services support pausing (e.g., Azure SQL Database serverless supports pausing, but not all offerings do). |

You have correctly selected 2.

The correct answer is: Platform as a service (PaaS) database offerings in Azure require less setup and configuration effort than infrastructure as a service (IaaS) database offerings. → YES, Platform as a service (PaaS) database offerings in Azure provide end users with the ability to control and update the operating system version. → NO, All relation and non-relational platform as a service (PaaS) database offerings in Azure can be paused to reduce costs. → NO

Question **27**

Incorrect

Mark 0.00 out of 1.00

Select the answer that correctly completes the sentence.

A database administrator is responsible for creating visuals and charts that help a company make informed decisions.

Your answer is incorrect.

The correct answer is:
Select the answer that correctly completes the sentence.

[A data analyst] is responsible for creating visuals and charts that help a company make informed decisions.

Question **28**

Incorrect

Mark 0.00 out of 1.00

You have an e-commerce application that reads and writes data to an Azure SQL database.
Which type of processing does the application use?

- a. stream processing
- b. batch processing
- c. Online Transaction Processing (OLTP)
- d. Online Analytical Processing (OLAP)

Your answer is incorrect.

## Correct Answer:

**Online Transaction Processing (OLTP)**

## Reason:

**OLTP** systems are designed for:

- Handling **frequent, real-time insert, update, and delete** operations,
- Supporting **day-to-day business operations** like order placement, inventory updates, or customer transactions.

An **e-commerce application** typically performs many **small, fast transactions**, making **OLTP** the correct model.

The correct answer is: Online Transaction Processing (OLTP)

Question **29**

Correct

Mark 1.00 out of 1.00

You have a SQL query that combines customer data and order data. The query includes calculated columns.
You need to create a database object that would allow other users to rerun the same SQL query.
What should you create?

- a. a view
- b. a table
- c. an index
- d. a scalar function

Your answer is correct.

The correct answer is: a view

Question **30**

Correct

Mark 1.00 out of 1.00

Which language is used to define queries in Azure Synapse Data Explorer?

- a. KQL
- b. SQL
- c. PowerShell
- d. Bash

Your answer is correct.

The correct answer is: KQL

Question **31**

Correct

Mark 1.00 out of 1.00

What is the responsibility of a database administrator?

- a. managing data integration processes
- b. applying data cleaning routines
- c. storing backup copies of data
- d. creating data visualizations

Your answer is correct.

The correct answer is:
storing backup copies of data

Question **32**

Correct

Mark 1.00 out of 1.00

Relational databases are commonly used to store and query [ structured data. ]

Your answer is correct.

The correct answer is:
Relational databases are commonly used to store and query [ structured data.]

Question **33**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

An elastic pool enables multiple Azure SQL databases to share resources such as memory, data storage space, and processing power.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

[An elastic pool] enables multiple Azure SQL databases to share resources such as memory, data storage space, and processing power.

Question **34**

Correct

Mark 1.00 out of 1.00

DRAG DROP -
Match the types of workloads to the appropriate scenarios.
To answer, drag the appropriate workload type from the column on the left to its scenario on the right. Each workload type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Batch — Analyzing historical data containing web traffic collected during the previous year.

Batch — Classifying images that were uploaded last month.

Streaming — Tracking in real time how many people are currently using a website.

Your answer is correct.

The correct answer is:
DRAG DROP -
Match the types of workloads to the appropriate scenarios.
To answer, drag the appropriate workload type from the column on the left to its scenario on the right. Each workload type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

[Batch ] Analyzing historical data containing web traffic collected during the previous year.

[Batch ]Classifying images that were uploaded last month.

[Streaming]Tracking in real time how many people are currently using a website.

Question **35**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Backing up and restoring data are responsibilities of a A database administrator .

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

Backing up and restoring data are responsibilities of a [A database administrator].

Question **36**

Correct

Mark 1.00 out of 1.00

A retail point of sale (POS) system is an example of which type of solution?

a. extract, transform, and load (ETL)

b. online analytical processing (OLAP)

c. online transaction processing (OLTP)

d. a data warehouse

Your answer is correct.

The correct answer is: online transaction processing (OLTP)

Question **37**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Social media users and followers in a graph database is an example of semi-structured data.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

Social media users and followers in a graph database is an example of [semi-structured data.]

Question **38**

Correct

Mark 1.00 out of 1.00

When you create an Azure SQL database, which account can always connect to the database?

a. the server admin login account of the logical server

b. the Azure Active Directory (Azure AD) administrator account

c. the sa account

d. the Azure Active Directory (Azure AD) account that created the database

Your answer is correct.

The correct answer is:
the server admin login account of the logical server

Question **39**

Partially correct

Mark 2.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Azure SQL Database includes a managed backup service. | YES |
| Azure SQL Database has built-in high availability. | NO |
| Azure SQL Database can use Azure Defender. | YES |

Your answer is partially correct.

| Statement | Yes/No | Reason |
|---|---|---|
| **Azure SQL Database includes a managed backup service.** | **Yes** | Azure SQL Database automatically performs **full, differential, and transaction log backups**, with **no manual intervention required**. |
| **Azure SQL Database has built-in high availability.** | **Yes** | High availability is **built-in by default**, using **replication and failover mechanisms** within the Azure platform. |
| **Azure SQL Database can use Azure Defender.** | **Yes** | Azure SQL Database can be integrated with **Microsoft Defender for SQL**, providing **threat protection and vulnerability assessments**. |

You have correctly selected 2.

The correct answer is: Azure SQL Database includes a managed backup service. → YES, Azure SQL Database has built-in high availability. → YES, Azure SQL Database can use Azure Defender. → YES

Question **40**

Correct

Mark 1.00 out of 1.00

What is a characteristic of a non-relational database?

a. full support for Transact-SQL

b. a fixed schema

c. self-describing entities

Your answer is correct.

The correct answer is: self-describing entities

Question **41**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| If you have a platform as a service (PaaS) database in Azure, you are responsible for applying operating system updates. | NO |
| If you have a platform as a service (PaaS) database in Azure, backups are performed automatically. | YES |
| If you have a platform as a service (PaaS) database in Azure, you are responsible for installation of the database engine. | NO |

Your answer is correct.

The correct answer is: If you have a platform as a service (PaaS) database in Azure, you are responsible for applying operating system updates. → NO, If you have a platform as a service (PaaS) database in Azure, backups are performed automatically. → YES, If you have a platform as a service (PaaS) database in Azure, you are responsible for installation of the database engine. → NO

Question **42**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point

Stream processing has access to the most recent data received or data within a rolling time window. Yes

Batch processing must occur immediately and have latency in the order of seconds or milliseconds. No

Stream processing is used for simple response functions, aggregates, or calculations such as rolling averages Yes .

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point

Stream processing has access to the most recent data received or data within a rolling time window. [Yes]
Batch processing must occur immediately and have latency in the order of seconds or milliseconds. [No]
Stream processing is used for simple response functions, aggregates, or calculations such as rolling averages [Yes].

Question **43**

Correct

Mark 1.00 out of 1.00

When can you use an Azure Resource Manager template?

to automate the creation of an interdependent group of Azure resources in a repeatable way

to apply Azure policies for multi-tenant deployments

to provision Azure subscriptions

- a. to automate the creation of an interdependent group of Azure resources in a repeatable way
- b. to apply Azure policies for multi-tenant deployments
- c. to control which services and feature administrators and developers can deploy from the Azure portal
- d. to provision Azure subscriptions

Your answer is correct.

The correct answer is:
to automate the creation of an interdependent group of Azure resources in a repeatable way

Question **44**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

In a relational database, the schema is defined before data is written to the table. Yes

Key/value data stores are highly optimized for complex queries that contain multiple joins. No

A relational database is commonly used to store and query structured data. Yes

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

In a relational database, the schema is defined before data is written to the table. [Yes]
Key/value data stores are highly optimized for complex queries that contain multiple joins. [No]
A relational database is commonly used to store and query structured data. [Yes]

Question **45**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

To add a new attribute to an entity in structured data, you need to

update all the records in the entity table to include the new column.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

To add a new attribute to an entity in structured data, you need to [update all the records in the entity table to include the new column.]

Question **46**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

Azure Databricks can consume data from Azure Event Hubs

Azure Databricks can consume data from Azure Cosmos DB

| | |
|---|---|
| Azure Databricks can consume data from Azure SQL Database | YES |
| Azure Databricks can consume data from Azure Event Hubs | YES |
| Azure Databricks can consume data from Azure Cosmos DB | YES |

Your answer is correct.

The correct answer is:
Azure Databricks can consume data from Azure SQL Database → YES,

Azure Databricks can consume data from Azure Event Hubs → YES,

Azure Databricks can consume data from Azure Cosmos DB → YES

Question **47**

Correct

Mark 3.00 out of 3.00

DRAG DROP -
Match the types of analytics that can be used to answer the business questions.
To answer, drag the appropriate analytics type from the column Each analytics type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Why did sales increase last month? Diagnostic

Provide specific guidance on how to allocate current budget to buy different inventory items. Prescriptive

Detect images of the company's logo included in PDF documents Cognitive

Select and Place:

Descriptive Predictive

Your answer is correct.

The correct answer is:
DRAG DROP -
Match the types of analytics that can be used to answer the business questions.
To answer, drag the appropriate analytics type from the column Each analytics type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Why did sales increase last month? [ Diagnostic ]

Provide specific guidance on how to allocate current budget to buy different inventory items.[ Prescriptive ]

Detect images of the company's logo included in PDF documents [Cognitive ]

Select and Place:

Question **48**

Correct

Mark 2.00 out of 2.00

Which two Azure services can be used to provision Apache Spark clusters? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [ ] a. Azure Time Series Insights
- [x] b. Azure Databricks
- [ ] c. Azure Log Analytics
- [x] d. Azure HDInsight

Your answer is correct.

The correct answers are:
Azure HDInsight,

Azure Databricks

Question **49**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

A Microsoft Power BI [paginated report] enables users to create highly formatted, fixed-layout documents optimized for printing and archiving.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

A Microsoft Power BI [paginated report] enables users to create highly formatted, fixed-layout documents optimized for printing and archiving.

Question **50**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

In Azure Data Factory, you can use [a control flow]
to orchestrate pipeline activities that depend on the output of other pipeline activities.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

In Azure Data Factory, you can use[a control flow]
to orchestrate pipeline activities that depend on the output of other pipeline activities.

Question **51**

Correct

Mark 2.00 out of 2.00

You have an Azure SQL database that contains tables for customers and orders.

You need to create a query that will return the orders placed by each customer.

Which two Transact-SQL statements should you include in the query? Each correct answer presents part of the solution.

NOTE: Each correct selection is worth one point.

- [ ] a. EXISTS
- [ ] b. INDEX
- [x] c. SELECT
- [ ] d. ORDERBY
- [x] e. JOIN

Your answer is correct.

The correct answers are: JOIN, SELECT

Question **52**

Correct

Mark 1.00 out of 1.00

Which database transaction property ensures that individual transactions are executed only once and either succeed in their entirety or roll back?

- a. durability
- b. atomicity
- c. isolation
- d. consistency

Your answer is correct.

The correct answer is: atomicity

Question **53**

Correct

Mark 1.00 out of 1.00

You plan to deploy an app. The app requires a nonrelational data service that will provide latency guarantees of less than 10-ms for reads and writes.
What should you include in the solution?

- a. Azure Files
- b. Azure Table storage
- c. Azure Blob storage
- d. Azure Cosmos DB ✓

Your answer is correct.

The correct answer is: Azure Cosmos DB

Question **54**

Correct

Mark 1.00 out of 1.00

What can be used with native notebook support to query and visualize data by using a web-based interface?

- a. Azure Databricks ✓
- b. Microsoft Power BI
- c. pgAdmin

Your answer is correct.

The correct answer is: Azure Databricks

Question **55**

Correct

Mark 1.00 out of 1.00

A bar chart showing year-to-date sales by region is an example of which type of analytics?

- a. prescriptive
- b. predictive
- c. descriptive ✓
- d. diagnostic

Your answer is correct.

The correct answer is: descriptive

Question **56**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

The JSON file format uses a hierarchical document schema to define data entities that have multiple attributes.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

The [JSON] file format uses a hierarchical document schema to define data entities that have multiple attributes.

Question **57**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

XML-formatted files are examples of semi-structured data.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

XML-formatted files are examples of [semi-structured data.]

Question **58**

Correct

Mark 1.00 out of 1.00

Which storage solution supports role-based access control (RBAC) at the file and folder level?

- a. Azure Queue storage
- b. Azure Data Lake Storage
- c. Azure Blob storage
- d. Azure Disk Storage

Your answer is correct.

The correct answer is:
Azure Data Lake Storage

Question **59**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

A data analyst is responsible for identifying which business rules must be applied to the data of a company.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

[A data analyst ] is responsible for identifying which business rules must be applied to the data of a company.

Question **60**

Incorrect

Mark 0.00 out of 1.00

You are writing a set of SQL queries that administrators will use to troubleshoot an Azure SQL database.
You need to embed documents and query results into a SQL notebook.
What should you use?

- a. Azure PowerShell
- b. Azure CLI
- c. Azure Data Studio
- d. Microsoft SQL Server Management Studio (SSMS) ⊗

Your answer is incorrect.

**Correct Answer:**

**Azure Data Studio**

**Reason:**

**Azure Data Studio** supports:

- **SQL notebooks**, which allow you to combine **text (Markdown)**, **SQL code**, and **query results** in a single interactive document,
- Ideal for documentation, analysis, and troubleshooting workflows.

It is the **only tool listed** that supports SQL **notebooks natively**.

.

The correct answer is:
Azure Data Studio

Question **61**

Correct

Mark 1.00 out of 1.00

You have an application that runs on Windows and requires access to a mapped drive.
Which Azure service should you use?

- a. Azure Table storage
- b. Azure Files ✓
- c. Azure Cosmos DB
- d. Azure Blob storage

Your answer is correct.

The correct answer is: Azure Files

Question **62**

Correct

Mark 1.00 out of 1.00

Which clause should you use in a SELECT statement to combine rows in one table with rows in another table?

- a. JOIN
- b. VALUES
- c. KEY
- d. SET

Your answer is correct.

The correct answer is:
JOIN

Question **63**

Correct

Mark 3.00 out of 3.00

Match the Azure Cosmos DB APIs to the appropriate data structures.

To answer, drag the appropriate API from the column on the left to its data structure on the right. Each API may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

| | |
|---|---|
| Azure Cosmos DB for Apache Gremlin API | Graph Data |
| Azure Cosmos DB for NoSQL API | JSON Documents |
| Azure Cosmos DB for Table API | Key/Value Data |
| Azure Cosmos DB for Apache Cassandra API | |

Your answer is correct.

The correct answer is:
Match the Azure Cosmos DB APIs to the appropriate data structures.

To answer, drag the appropriate API from the column on the left to its data structure on the right. Each API may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

[Azure Cosmos DB for Apache Gremlin API] Graph Data

[Azure Cosmos DB for NoSQL API] JSON Documents

[Azure Cosmos DB for Table API] Key/Value Data

Question **64**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Stream processing can process all the data in a dataset at once. No

Batch processing occurs immediately, with latency in the order of seconds or milliseconds. No

Stream processing is used for simple response functions, aggregates, or calculations, such as rolling averages. Yes

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Stream processing can process all the data in a dataset at once. [No]
Batch processing occurs immediately, with latency in the order of seconds or milliseconds.[No]
Stream processing is used for simple response functions, aggregates, or calculations, such as rolling averages. [Yes]

Question **65**

Correct

Mark 1.00 out of 1.00

You have an Azure SQL database that you access directly from the Internet.
You recently changed the public IP address of your computer.
After changing the IP address, you can no longer access the database. You can connect to other resources in Azure.
What is a possible cause of the issue?

- a. role-based access control (RBAC
- b. Domain Name Servic (DNS)
- c. Dynamic Host Configuration Protocol (DHCP)
- d. a database-level firewall

Your answer is correct.

The correct answer is:
a database-level firewall

Question **66**

Correct

Mark 1.00 out of 1.00

You have a quality assurance application that reads data from a data warehouse.
Which type of processing does the application use?

- a. stream processing
- b. Online Analytical Processing (OLAP)
- c. batch processing
- d. Online Transaction Processing (OLTP)

Your answer is correct.

The correct answer is:
Online Analytical Processing (OLAP)

Question **67**

Correct

Mark 1.00 out of 1.00

You need to create an Azure Storage account.
Data in the account must replicate outside the Azure region automatically.
Which two types of replication can you use for the storage account? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [ ] a. zone-redundant storage (ZRS)
- [x] b. read-access geo-redundant storage (RA-GRS)
- [x] c. geo-redundant storage (GRS)
- [ ] d. locally-redundant storage (LRS)

Your answer is correct.

The correct answers are: read-access geo-redundant storage (RA-GRS), geo-redundant storage (GRS)

Question **68**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Stream processing is used for simple analytics with short latency.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

Stream processing is used for [simple analytics with short] latency.

Question **69**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

A visualization that shows a university's current student enrollment versus the maximum capacity is an example of descriptive analytics.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

A visualization that shows a university's current student enrollment versus the maximum capacity is an example of [descriptive]analytics.

Question **70**

Correct

Mark 3.00 out of 3.00

Match the data storage options to the appropriate types of data.

To answer, drag the appropriate data storage option from the column on the left to its data type on the right. Each data storage option may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

| | |
|---|---|
| Azure Blob Storage | Audio files |
| Azure Cosmos DB for NoSQL | JSON documents |
| Azure SQL Database | Tabular datasets |

Azure Data Factory

Your answer is correct.

The correct answer is:
Match the data storage options to the appropriate types of data.

To answer, drag the appropriate data storage option from the column on the left to its data type on the right. Each data storage option may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

[Azure Blob Storage] Audio files

[Azure Cosmos DB for NoSQL] JSON documents

[Azure SQL Database] Tabular datasets

Question **71**

Correct

Mark 3.00 out of 3.00

DRAG DROP -
Match the security components to the appropriate scenarios.
To answer, drag the appropriate component from the column. Each component may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Prevent access to an Azure SQL database from another network. Firewall

Support Azure Active Directory (Azure AD) sign-ins to an Azure SQL database. Authentication

Ensure that sensitive data never appears as plain text in an Azure SQL database. Encryption

Select and Place:

Your answer is correct.

The correct answer is:
DRAG DROP -
Match the security components to the appropriate scenarios.
To answer, drag the appropriate component from the column. Each component may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Prevent access to an Azure SQL database from another network.[Firewall]

Support Azure Active Directory (Azure AD) sign-ins to an Azure SQL database.[Authentication]

Ensure that sensitive data never appears as plain text in an Azure SQL database.[Encryption]
Select and Place:

Question **72**

Correct

Mark 1.00 out of 1.00

Your company needs to design a database that shows how changes in network traffic in one area of a network affect network traffic in other areas of the network.
Which type of data store should you use?

- a. document
- b. columnar
- c. key/value
- d. graph

Your answer is correct.

The correct answer is: graph

Question **73**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

The Azure Cosmos DB [for NoSQL] ✓

API enables the use of SELECT statements to retrieve documents from Azure Cosmos DB.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

The Azure Cosmos DB [for NoSQL]

API enables the use of SELECT statements to retrieve documents from Azure Cosmos DB.

Question **74**

Correct

Mark 2.00 out of 2.00

Match the Azure Data Lake Storage Gen2 terms to the appropriate levels in the hierarchy.

To answer, drag the appropriate term from the column on the left to its level on the right. Each term may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

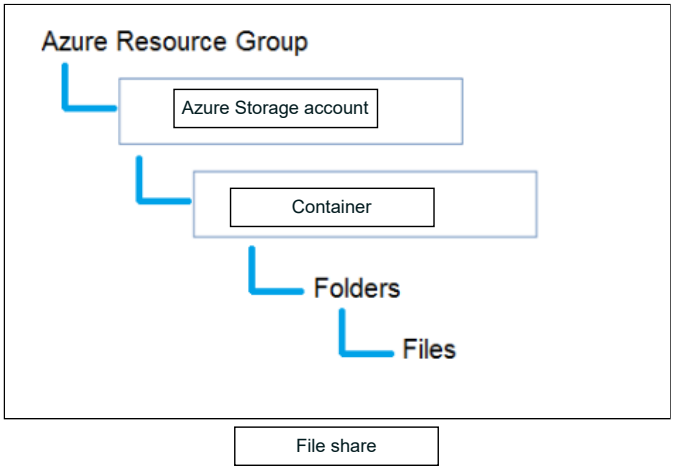


Your answer is correct.

Question **75**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

When provisioning an Azure Cosmos DB [account] you need to specify which type of API you will use.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.
When provisioning an Azure Cosmos DB [account ] you need to specify which type of API you will use.

Question **76**

Correct

Mark 1.00 out of 1.00

[Microsoft Power BI Desktop] provides a full range of Power BI data modeling and report editing features.

Your answer is correct.

The correct answer is:
[Microsoft Power BI Desktop] provides a full range of Power BI data modeling and report editing features.

Question **77**

Partially correct

Mark 1.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| A pipeline is a representation of a data structure within Azure Data Factory | NO |
| Azure Data Factory pipelines can execute other pipelines | NO |
| A processing step within an Azure Data Factory pipeline is an activity | NO |

Your answer is partially correct.

You have correctly selected 1.

The correct answer is:
A pipeline is a representation of a data structure within Azure Data Factory → NO,

Azure Data Factory pipelines can execute other pipelines → YES,

A processing step within an Azure Data Factory pipeline is an activity → YES

Question **78**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

The CSV format is used to store delimited data, separates data fields by using commas, and terminates rows by using a carriage return.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

The [CSV] format is used to store delimited data, separates data fields by using commas, and terminates rows by using a carriage return.

Question **79**

Correct

Mark 1.00 out of 1.00

A bank needs to ensure that after an account transfer transaction completes, the revised account balances persists even if the database system hosting the transaction becomes temporarily unavailable.

Of which ACID semantic is this an example?

- a. isolation
- b. durability
- c. consistency
- d. atomicity

Your answer is correct.

The correct answer is: durability

Question **80**

Correct

Mark 1.00 out of 1.00

What is a primary characteristic of a relational database?

- a. a flexible data structure
- b. a large amount of duplicate data
- c. a lack of dependencies between tables
- d. data is queried and manipulated by using a variant of the SQL language

Your answer is correct.

The correct answer is: data is queried and manipulated by using a variant of the SQL language

Question **81**

Partially correct

Mark 2.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | | |
|---|---|---|
| A Microsoft Power BI dashboard is associated with a single workspace. | NO | ✗ |
| A Microsoft Power BI dashboard can only display visualizations from a single dataset. | NO | ✓ |
| A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook. | YES | ✓ |

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: A Microsoft Power BI dashboard is associated with a single workspace. → YES, A Microsoft Power BI dashboard can only display visualizations from a single dataset. → NO, A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook. → YES

Question **82**

Correct

Mark 1.00 out of 1.00

Which type of database contains nodes and edges?

- a. graph ✓
- b. time series
- c. columnar
- d. key/value

Your answer is correct.

The correct answer is:
graph

Question **83**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

A key/value data store is optimized for simple lookups.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

A key/value data store is optimized for [simple lookups. ]

Question **84**

Incorrect

Mark 0.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

By default, each Azure SQL database is protected by Azure Firewall.

Your answer is incorrect.

**Correct Answer:**

A server-level firewall

**Explanation:**

By default, each **Azure SQL Database** is protected by a **server-level firewall** that:

- Restricts access to the **logical server** hosting the database
- Requires you to explicitly allow **IP address ranges** or **Azure services**
- Acts as the **first layer of defense** before connection attempts reach the database

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

By default, each Azure SQL database is protected by [a server-level firewall.]

Question **85**

Correct

Mark 2.00 out of 2.00

DRAG DROP -

Select the datastore services to the appropriate descriptions.

NOTE: Each correct match is worth one point.

| | |
|---|---|
| Azure Cosmos DB | Enables the use of SQL queries against data stored in JSON documents |
| Azure Files | Enables users to access data by using the Server Message Block (SMB) version 3 protocol |

Your answer is correct.

The correct answer is:
DRAG DROP -
Select the datastore services to the appropriate descriptions.
NOTE: Each correct match is worth one point.

[Azure Cosmos DB ] Enables the use of SQL queries against data stored in JSON documents

[Azure Files ] Enables users to access data by using the Server Message Block (SMB) version 3 protocol

Question **86**

Correct

Mark 1.00 out of 1.00

You have a transactional application that stores data in an Azure SQL managed instance.
When should you implement a read-only database replica?

a. You need to improve the recovery point objective (RPO).

b. You need to generate reports without affecting the transactional workload.

c. You need to audit the transactional application.

d. You need to implement high availability in the event of a regional outage.

Your answer is correct.

The correct answer is:
You need to generate reports without affecting the transactional workload.

Question **87**

Correct

Mark 3.00 out of 3.00

You need to query a table named Products in an Azure SQL database.
Which three requirements must be met to query the table from the internet? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point

- [x] a. Your IP address must be allowed to connect to the database.
- [ ] b. You must be assigned the Contributor role for the resource group that contains the database.
- [x] c. You must have a user in the database.
- [ ] d. You must be assigned the Reader role for the resource group that contains the database.
- [x] e. You must have SELECT access to the Products table.

Your answer is correct.

The correct answers are:
You must have SELECT access to the Products table.,

You must have a user in the database.,

Your IP address must be allowed to connect to the database.

Question **88**

Correct

Mark 3.00 out of 3.00

DRAG DROP -
Match the Azure SQL services to the appropriate use cases.
To answer, drag the appropriate service from the column on the left to its use case on the right. Each service may be used once, more than once, or not at all.
NOTE: Each correct match its worth one point.

| | |
|---|---|
| Azure SQL Database single database | A serverless configuration |
| Azure SQL Managed instance | A database that has the highest compatibility with on-premises Microsoft SQL Server |
| Azure SQL Database elastic pool | A low-cost, low-maintenance database |

Select and Place:

Your answer is correct.

The correct answer is:
DRAG DROP -
Match the Azure SQL services to the appropriate use cases.
To answer, drag the appropriate service from the column on the left to its use case on the right. Each service may be used once, more than once, or not at all.
NOTE: Each correct match its worth one point.

[Azure SQL Database single database ] A serverless configuration

[Azure SQL Managed instance] A database that has the highest compatibility with on-premises Microsoft SQL Server

[Azure SQL Database elastic pool ] A low-cost, low-maintenance database

Select and Place:

Question **89**

Correct

Mark 1.00 out of 1.00

Online analytical processing (OAP) systems are particularly suitable for analytical workloads because data is pre-aggregated

Your answer is correct.

The correct answer is:
Online analytical processing (OAP) systems are particularly suitable for analytical workloads because data is [pre-aggregated]

Question **90**

Correct

Mark 4.00 out of 4.00

Match the ACID terms to the appropriate descriptions.

To answer, drag the appropriate term from the column on the left to its description on the right. Each term may be used once, more than once or not at all.

| | |
|---|---|
| Isolation | Concurrent transactions cannot interfere with one another and must result in a consistent database state. |
| Atomicity | Each transaction is treated as a single unit that succeeds completely or fails completely. |
| Consistency | Transactions can only take the data in a database from one valid state to another. |
| Durability | When a transaction is committed, it will remain committed. |

Your answer is correct.

The correct answer is:
Match the ACID terms to the appropriate descriptions.

To answer, drag the appropriate term from the column on the left to its description on the right. Each term may be used once, more than once or not at all.

[Isolation] Concurrent transactions cannot interfere with one another and must result in a consistent database state.
[Atomicity] Each transaction is treated as a single unit that succeeds completely or fails completely.
[Consistency] Transactions can only take the data in a database from one valid state to another.
[Durability] When a transaction is committed, it will remain committed.

Question **91**

Incorrect

Mark 0.00 out of 1.00

You need to design and model a database by using a graphical tool that supports project-oriented offline database development.
What should you use?

- a. Azure Databricks
- b. Azure Data Studio
- c. Microsoft SQL Server Management Studio (SSMS)
- d. Microsoft SQL Server Data Tools (SSDT)

Your answer is incorrect.

**Correct Answer:**

Microsoft SQL Server Data Tools (SSDT)

**Explanation:**

**Microsoft SQL Server Data Tools (SSDT)** is designed for:

- **Database design and modeling**
- **Project-oriented, offline development**
- Creating **schemas, tables, views, stored procedures** in a disconnected project format
- Source control integration for **versioning** and team collaboration

It's ideal for building and maintaining **SQL Server** and **Azure SQL databases** in a **development lifecycle workflow**.

The correct answer is:
Microsoft SQL Server Data Tools (SSDT)

Question **92**

Partially correct

Mark 0.20 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

To configure an Azure Storage account to support access control lists

that have object-level permissions, enable the hierarchical namespace.

Your answer is partially correct.

You have correctly selected 1.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

To configure an Azure Storage account to support access control lists

that have object-level permissions, [enable the hierarchical namespace.]

Question **93**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

In an infrastructure as a service (IaaS) instance of Microsoft SQL Server on Azure, you manage the virtual machine that hosts SQL Server.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

In an infrastructure as a service (IaaS) instance of Microsoft SQL Server on Azure, you manage the [virtual machine] that hosts SQL Server.

Question **94**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

The WHERE clause can be used in Data Manipulation Language (DML) statements to specify the criteria that rows must match.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

The [WHERE] clause can be used in Data Manipulation Language (DML) statements to specify the criteria that rows must match.

Question **95**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

In a relational database, primary key the uniquely identifies each row in a table.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

In a relational database, [primary key ] the uniquely identifies each row in a table.

Question **96**

Correct

Mark 3.00 out of 3.00

Which three objects can be added to a Microsoft Power BI dashboard? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [x] a. a visualization from a report
- [x] b. a report page
- [ ] c. a Microsoft PowerPoint slide
- [x] d. a text box
- [ ] e. a dataflow

Your answer is correct.

The correct answers are:
a report page,

a visualization from a report,

a text box

Question **97**

Correct

Mark 3.00 out of 3.00

Match the database normalization terms to the appropriate descriptions.

To answer, drag the appropriate term from the column on the left to its description on the right. Each term may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

| | |
|---|---|
| Attributes | Get separated into individual columns |
| Entities | Get separated into individual tables |
| Foreign keys | Link two related entities |

Attributes | Entities | Foreign keys | Primary keys | Schemas

Your answer is correct.

The correct answer is:
Match the database normalization terms to the appropriate descriptions.

To answer, drag the appropriate term from the column on the left to its description on the right. Each term may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

[Attributes] Get separated into individual columns
[Entities] Get separated into individual tables
[Foreign keys] Link two related entities

Question **98**

Correct

Mark 3.00 out of 3.00

DRAG DROP -

Match the types of data to the appropriate Azure data services.

To answer, drag the appropriate data type from the column on the left to its service on the right. Each data type may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

| | | |
|---|---|---|
| Image files | ✓ | Azure Blob storage |
| Relationships between employees | ✓ | Azure Cosmos DB Gremlin API |
| Key/value pairs | ✓ | Azure Table storage |

Select and Place:

Your answer is correct.

The correct answer is:
DRAG DROP -

Match the types of data to the appropriate Azure data services.

To answer, drag the appropriate data type from the column on the left to its service on the right. Each data type may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

[Image files] Azure Blob storage

[Relationships between employees] Azure Cosmos DB Gremlin API

[Key/value pairs] Azure Table storage

Select and Place:

Question **99**

Correct

Mark 1.00 out of 1.00

You need to accept and store large quantities of IoT streaming data.

Which database solution should you use?

- a. Azure Cache for Redis
- b. Azure Cosmos DB ✓
- c. Azure Database for MariaDB
- d. Azure Database for PostgreSQL

Your answer is correct.

The correct answer is: Azure Cosmos DB

Question **100**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Each Azure SQL managed instance supports multiple databases that can be accessed by using cross-database queries. | YES |
| Elastic pools are used to share resources across multiple instances of SQL Server on Azure Virtual Machines. | NO |
| Azure SQL Database is fully compatible with both on-premises physical instances and virtualized instances of Microsoft SQL Server. | NO |

Your answer is correct.

The correct answer is:
Each Azure SQL managed instance supports multiple databases that can be accessed by using cross-database queries. → YES,

Elastic pools are used to share resources across multiple instances of SQL Server on Azure Virtual Machines. → NO,

Azure SQL Database is fully compatible with both on-premises physical instances and virtualized instances of Microsoft SQL Server. → NO

Question **101**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

An extract load, and transform (ELT) process requires a target data store powerful enough to transform data.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

An extract load, and transform (ELT) process requires [a target data store powerful enough to transform data.]

Question **102**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Platform as a service (PaaS) database offerings in Azure provide built-in high availability. | YES |
| Platform as a service (PaaS) database offerings in Azure provide configurable scaling options. | YES |
| Platform as a service (PaaS) database offerings in Azure reduce the administrative overhead for managing hardware. | YES |

Your answer is correct.

The correct answer is:
Platform as a service (PaaS) database offerings in Azure provide built-in high availability. → YES, Platform as a service (PaaS) database offerings in Azure provide configurable scaling options. → YES,

Platform as a service (PaaS) database offerings in Azure reduce the administrative overhead for managing hardware. → YES

Question **103**

Correct

Mark 2.00 out of 2.00

Which two activities can be performed entirely by using the Microsoft Power BI service without relying on Power BI Desktop? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [ ] a. data acquisition and preparation
- [x] b. report and dashboard creation
- [x] c. report sharing and distribution
- [ ] d. data modeling

Your answer is correct.

The correct answers are:
report and dashboard creation,

report sharing and distribution

Question **104**

Correct

Mark 1.00 out of 1.00

Your company is designing a database that will contain session data for a website. The data will include notifications, personalization attributes, and products that are added to a shopping cart.
Which type of data store will provide the lowest latency to retrieve the data?

- a. document
- b. columnar
- c. graph
- d. key/value ✓

Your answer is correct.

The correct answer is: key/value

Question **105**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | | |
|---|---|---|
| Azure Table storage supports multiple read replicas. | YES | ✓ |
| Azure Table storage supports multiple write regions. | NO | ✓ |
| The Azure Cosmos DB Table API supports multiple read replicas. | YES | ✓ |
| The Azure Cosmos DB Table API supports multiple write regions. | YES | ✓ |

Your answer is correct.

The correct answer is: Azure Table storage supports multiple read replicas. → YES, Azure Table storage supports multiple write regions. → NO, The Azure Cosmos DB Table API supports multiple read replicas. → YES, The Azure Cosmos DB Table API supports multiple write regions. → YES

Question **106**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completed the sentence.

Data in a data warehouse can be aggregated and loaded into an online analytical processing (OLAP) model, which is also known as a cube.

Your answer is correct.

The correct answer is:
Select the answer that correctly completed the sentence.

Data in a data warehouse can be aggregated and loaded into an online analytical processing (OLAP) model, which is also known as a [cube.]

Question **107**

Correct

Mark 1.00 out of 1.00

Which Azure Storage service implements the key/value model?

a. Azure Blob

b. Azure Queue

c. Azure Table

d. Azure Files

Your answer is correct.

The correct answer is: Azure Table

Question **108**

Correct

Mark 1.00 out of 1.00

You have data saved in the following format.

```
<Person Age= "23">
     <FirstName>John</FirstName>
     <LastName>Smith</LastName>
     <Hobbies>
          <Hobby Type= "Sports">Golf</Hobby>
          <Hobby Type= "Leisure">Reading</Hobby>
          <Hobby Type= "Leisure">Guitar</Hobby>
     </Hobbies>
</Person>
```

Which format was used?

- a. XML
- b. HTML
- c. YAML
- d. JSON

Your answer is correct.

The correct answer is: XML

Question **109**

Correct

Mark 1.00 out of 1.00

You have an Azure subscription that contains a Log Analytics workspace.

You need to onboard Microsoft Sentinel.

What should you do first?

- a. Correlate alerts into incidents.
- b. Create a hunting query.
- c. Connect to your security sources.
- d. Create a custom detection rule.

Your answer is correct.

The correct answer is:
Connect to your security sources.

Question **110**

Correct

Mark 1.00 out of 1.00

DRAG DROP -

Match the tools to the appropriate descriptions.

To answer, drag the appropriate tool from the column. Each tool may be used once, more than once, or not at all.

Microsoft SQL Server Management Studio (SSMS) A graphical tool for managing SQL Server or Azure SQL databases that supports access, configuration, management, and administration tasks.

Azure Data Studio A lightweight editor that can run on-demand SQL queries and view and sav results as text, JSON, or Microsoft Excel files.

Microsoft SQL Server Data Tools (SSDT) A development tool for building Azure SQL databases, Microsoft SQL Serve relational databases, SQL Server Analysis Services (SSAS) data models. SQL Server Integration Services (SSIS) packages, and SQL Server Reporting Services (SSBS) reports

Select and Place:

Your answer is correct.

The correct answer is:
DRAG DROP -
Match the tools to the appropriate descriptions.
To answer, drag the appropriate tool from the column. Each tool may be used once, more than once, or not at all.

[Microsoft SQL Server Management Studio (SSMS)] A graphical tool for managing SQL Server or Azure SQL databases that supports access, configuration, management, and administration tasks.

[Azure Data Studio ] A lightweight editor that can run on-demand SQL queries and view and sav results as text, JSON, or Microsoft Excel files.

[Microsoft SQL Server Data Tools (SSDT)] A development tool for building Azure SQL databases, Microsoft SQL Serve relational databases, SQL Server Analysis Services (SSAS) data models. SQL Server Integration Services (SSIS) packages, and SQL Server Reporting Services (SSBS) reports

Select and Place:

Question **111**

Correct

Mark 1.00 out of 1.00

A bank needs to ensure that a transaction involving debiting funds from a source account and crediting the same funds to a destination account must complete both actions. If either action falls to complete, the other action must fail.

Of which ACID semantic is this an example?

a. consistency

b. durability

c. atomicity

d. isolation

Your answer is correct.

The correct answer is: atomicity

Question **112**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

A job that calculates a rolling average temperature reading is an example of a streaming workload.
A job that calculates average revenue per product for the last month is an example of a streaming workload.

| | |
|---|---|
| A job that processes sales data once daily is an example of a batch workload. | YES |
| A job that calculates a rolling average temperature reading is an example of a streaming workload. | YES |
| A job that calculates average revenue per product for the last month is an example of a streaming workload. | NO |

Your answer is correct.

The correct answer is: A job that processes sales data once daily is an example of a batch workload. → YES,
A job that calculates a rolling average temperature reading is an example of a streaming workload. → YES,

A job that calculates average revenue per product for the last month is an example of a streaming workload. → NO

Question **113**

Incorrect

Mark 0.00 out of 2.00

Match the Azure services to the appropriate locations in the architecture.
To answer, drag the appropriate service from the column on the left to its location on the right. Each service may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

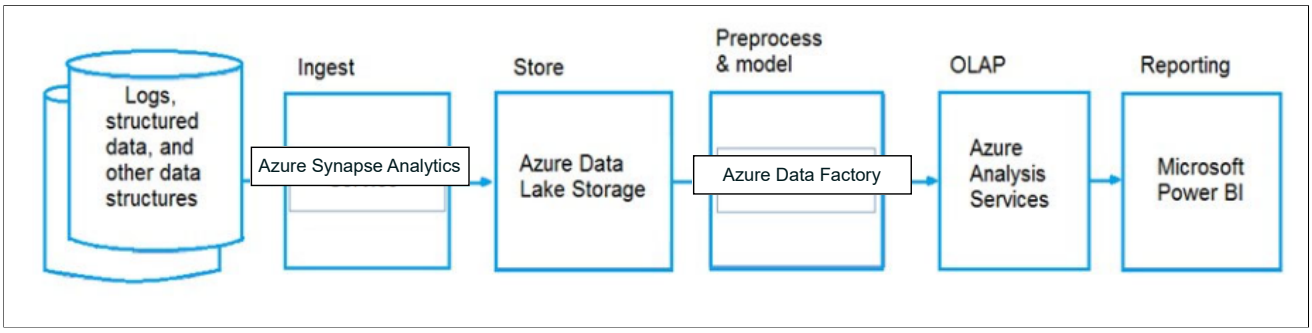


Azure Cognitive Search | Azure Data Catalog

Your answer is incorrect.

Question **114**

Correct

Mark 1.00 out of 1.00

You have the following data.

```
{
  "firstName": "John",
  "lastName": "Doe",
  "address":
  {
    "streetAddress": "1 Main St.",
    "city": "New York",
    "state": "NY",
    "postalCode": "10099"
  }
},
{
  "firstName": "Jane",
  "lastName": "Doe",
  "address":
  {
    "streetAddress": "1 Main St.",
    "city": "New York",
    "state": "NY",
    "postalCode": "10099"
  },
  "contact":
  [
    {
      "type": "email",
      "address": "jane.doe@contoso.com"
    }
  ]
}
```

Which type of data is this?

- a. semi-structured ✓
- b. structured
- c. unstructured

Your answer is correct.

The correct answer is: semi-structured

Question **115**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

A clustered index is an object associated with a table that sorts and stores the data rows in the table based on their key values.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

[A clustered index] is an object associated with a table that sorts and stores the data rows in the table based on their key values.

Question **116**

Partially correct

Mark 2.00 out of 3.00

DRAG DROP -

Match the types of analytics that can be used to answer the business questions.

To answer, drag the appropriate analytics type from the column .Each analytics type may be used once, more than once, or not at all.

Why did sales increase last month? Descriptive

How do I allocate my budget to buy different inventory items? Predictive

Which people are mentioned in a company's business documents? Cognitive

Select and Place:

Analytics Types

Diagnostic Prescriptive

Your answer is partially correct.

| **Business Question** | **Correct Analytics Type** |
|---|---|
| Why did sales increase last month? | **Diagnostic** |
| How do I allocate my budget to buy different inventory items? | **Prescriptive** |
| Which people are mentioned in a company's business documents? | **Cognitive** |

## Reasons:

- **Diagnostic Analytics** → Explains **why** something happened. Used for root cause analysis.
- **Prescriptive Analytics** → Recommends **what actions to take**, like optimizing budget allocation.
- **Cognitive Analytics** → Involves **AI-driven insights** such as understanding natural language in documents and identifying entities like names.

You have correctly selected 2.

The correct answer is:
DRAG DROP -

Match the types of analytics that can be used to answer the business questions.

To answer, drag the appropriate analytics type from the column .Each analytics type may be used once, more than once, or not at all.

Why did sales increase last month? [Diagnostic]

How do I allocate my budget to buy different inventory items?[Predictive]

Which people are mentioned in a company's business documents?[Cognitive]

Select and Place:

Analytics Types

Question **117**

Partially correct

Mark 2.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Azure Data Lake storage can be used as a staging area for ingested data before the data is converted into a format suitable for performing analytics. Yes

Azure Data Factory is an analytics engine that uses a massively parallel processing (MPP) architecture. No

Azure HDInsight enables you to build tabular models from multiple data sources to support online analytical processing (OLAP) queries. Yes

Your answer is partially correct.

You have correctly selected 2.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

Azure Data Lake storage can be used as a staging area for ingested data before the data is converted into a format suitable for performing analytics. [Yes]
Azure Data Factory is an analytics engine that uses a massively parallel processing (MPP) architecture. [No]
Azure HDInsight enables you to build tabular models from multiple data sources to support online analytical processing (OLAP) queries. [No]

Question **118**

Incorrect

Mark 0.00 out of 1.00

You need to store event log data that is semi-structured and received as the logs occur.
What should you use?

a. Azure Queue storage

b. Azure Table storage

c. Azure Files

Your answer is incorrect.

The correct answer is: Azure Table storage

Question **119**

Correct

Mark 2.00 out of 2.00

What are two characteristics of real-time data processing? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [ ] a. Data is processed periodically
- [x] b. Data is processed as it is created
- [ ] c. High latency is acceptable
- [x] d. Low latency is expected

Your answer is correct.

The correct answers are: Low latency is expected, Data is processed as it is created

Question **120**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

Relational data uses keys to enforce relationships between different tables.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

Relational data uses [keys] to enforce relationships between different tables.

Question **121**

Correct

Mark 1.00 out of 1.00

You have an inventory management database that contains the following table.

| ProductName | Quantity |
|---|---|
| Product1 | 100 |
| Product2 | 129 |
| Product3 | 176 |

Which statement should you use in a SQL query to change the inventory quantity of Product1 to 270?

- a. MERGE
- b. CREATE
- c. UPDATE
- d. INSERT

Your answer is correct.

The correct answer is:
UPDATE

Question **122**

Incorrect

Mark 0.00 out of 1.00

Which statement is an example of Data Manipulation Language (DML)?

- a. UPDATE
- b. CREATE
- c. REVOKE
- d. DISABLE

Your answer is incorrect.

**Correct Answer:** UPDATE

**Reason:**

**DML (Data Manipulation Language)** deals with operations that **modify the data** within existing tables:

- Examples include **INSERT**, **UPDATE**, **DELETE**, and **SELECT**.

**UPDATE** is used to **modify existing records**, making it a classic DML operation.

The correct answer is: UPDATE

Question **123**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

In Microsoft Power BI, a report is a collection of visualizations that appear together on one or more pages.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

In Microsoft Power BI, a [report] is a collection of visualizations that appear together on one or more pages.

Question **124**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

You can query a graph database in Azure Cosmos DB as nodes and edges by using the Gremlin language.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

You can query a graph database in Azure Cosmos DB [as nodes and edges by using the Gremlin language.]

Question **125**

Correct

Mark 1.00 out of 1.00

Which file format has defined names and data types for each column and uses compressed columnar storage?

a. Apache Parquet

b. CSV

c. JSON

d. HTML

Your answer is correct.

The correct answer is: Apache Parquet

Question **126**

Incorrect

Mark 0.00 out of 1.00

You have an on-premises Microsoft SQL Server database.
You need to migrate the database to the cloud. The solution must meet the following requirements:
* Minimize maintenance effort.
* Support the Database Mail and Service Broker features.
What should you include in the solution?

a. Azure SQL Managed instance

b. SQL Server on Azure virtual machines

c. an Azure SQL Database elastic pool

d. Azure SQL Database single database

Your answer is incorrect.

The correct answer is: Azure SQL Managed instance

Question **127**

Partially correct

Mark 2.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Azure DDoS Protection Standard protects against man-in-the-middle (MITM) attacks. | YES |
| Azure DDoS Protection Standard is enabled by default in an Azure subscription. | NO |
| Azure DDoS Protection Standard protects against protocol attacks. | YES |

Your answer is partially correct.

**Azure DDoS Protection Standard protects against man-in-the-middle (MITM) attacks.:No**
**Explanation:** MITM attacks target data interception, not volume-based traffic disruptions. DDoS protection doesn't prevent MITM attacks.

**Azure DDoS Protection Standard is enabled by default in an Azure subscription.: No**
**Explanation:** It is **not** enabled by default. You must **manually enable and configure** it on a virtual network.

**Azure DDoS Protection Standard protects against protocol attacks.: Yes**
**Explanation:** It does protect against protocol attacks (e.g., SYN floods, UDP floods) that attempt to exhaust system resources.

You have correctly selected 2.

The correct answer is: Azure DDoS Protection Standard protects against man-in-the-middle (MITM) attacks. → NO, Azure DDoS Protection Standard is enabled by default in an Azure subscription. → NO, Azure DDoS Protection Standard protects against protocol attacks. → YES

Question **128**

Correct

Mark 1.00 out of 1.00

What is a characteristic of non-relational data?

- a. no unique key values
- b. a well-defined schema
- c. self-describing entities ✓
- d. no indexes

Your answer is correct.

The correct answer is: self-describing entities

Question **129**

Incorrect

Mark 0.00 out of 1.00

Which database transaction property ensures that transactional changes to a database are preserved during unexpected operating system restarts?

- a. isolation ✗
- b. atomicity
- c. durability
- d. consistency

Your answer is incorrect.

The correct answer is: durability

Question **130**

Correct

Mark 1.00 out of 1.00

You have a data store that has the structure shown in the following exhibit.

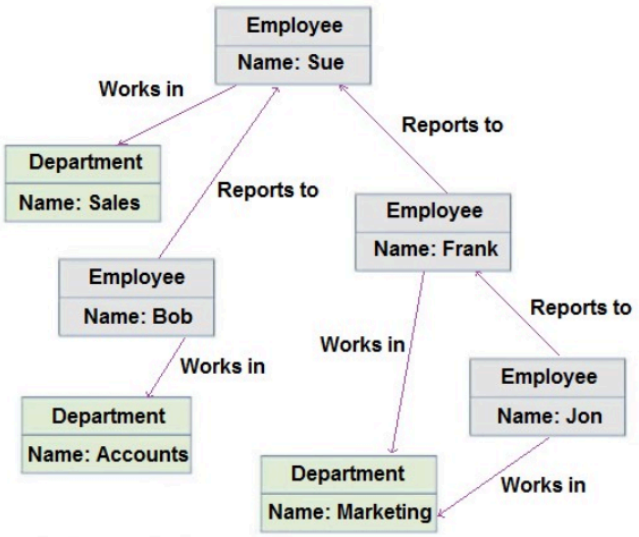


Which type of data store is this?

- a. graph ✓
- b. key/value
- c. object data
- d. time series

Your answer is correct.

The correct answer is: graph

Question **131**

Correct

Mark 1.00 out of 1.00

You need to implement an Azure platform as a service (PaaS) service that will host a relational database. The solution must support built-in autoscaling.

Which service should you implement?

- a. Azure SQL Managed Instance
- b. Azure SQL Database ✓
- c. SQL Server on Azure Virtual Machines
- d. Azure SQL Edge

Your answer is correct.

The correct answer is:
Azure SQL Database

Question **132**

Correct

Mark 3.00 out of 3.00

Match the data processing languages to the appropriate real-time analytics services.

To answer, drag the appropriate language from the column on the left to its service on the right. Each language may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

| | |
|---|---|
| KQL | Azure Data Explorer |
| SQL | Azure Stream Analytics |
| Python | Azure Synapse Analytics |
| JavaScript | |

Your answer is correct.

The correct answer is:
Match the data processing languages to the appropriate real-time analytics services.

To answer, drag the appropriate language from the column on the left to its service on the right. Each language may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

[KQL] Azure Data Explorer

[SQL] Azure Stream Analytics

[Python] Azure Synapse Analytics

Question **133**

Correct

Mark 1.00 out of 1.00

You need to recommend a data store service that meets the following requirements:
- Native SQL API access
- Configurable indexes

What should you recommend?

a. Azure Table storage

b. Azure Files

c. Azure Cosmos DB

d. Azure Blob storage

Your answer is correct.

The correct answer is: Azure Cosmos DB

Question **134**

Correct

Mark 1.00 out of 1.00

Which property of a transactional workload guarantees that each transaction is treated as a single unit that either succeeds completely or fails completely?

- a. atomicity
- b. isolation
- c. consistency
- d. durability

Your answer is correct.

The correct answer is:
atomicity

Question **135**

Correct

Mark 4.00 out of 4.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Azure Table storage within a single Azure Storage account supports multiple concurrent reads in different Azure regions. | YES |
| Azure Table storage within a single Azure Storage account supports multiple concurrent writes in different Azure regions. | NO |
| An Azure Cosmos DB account that uses the Table API supports multiple concurrent reads in different Azure regions. | YES |
| An Azure Cosmos DB account that uses the Table API supports multiple concurrent writes in different Azure regions. | YES |

Your answer is correct.

The correct answer is: Azure Table storage within a single Azure Storage account supports multiple concurrent reads in different Azure regions. → YES, Azure Table storage within a single Azure Storage account supports multiple concurrent writes in different Azure regions. → NO, An Azure Cosmos DB account that uses the Table API supports multiple concurrent reads in different Azure regions. → YES, An Azure Cosmos DB account that uses the Table API supports multiple concurrent writes in different Azure regions. → YES

Question **136**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Database administrators ✓ are responsible for the availability, performance, and optimization of databases.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

[Database administrators] are responsible for the availability, performance, and optimization of databases.

Question **137**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point

Batch processing is optimized for handling large datasets. Yes ✓

Stream processing is optimized for performing complex analysis of datasets. No ✓

The latency of batch processing is typically measured in seconds or milliseconds. No ✓

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point

Batch processing is optimized for handling large datasets. [Yes]
Stream processing is optimized for performing complex analysis of datasets. [No]
The latency of batch processing is typically measured in seconds or milliseconds. [No]

Question **138**

Correct

Mark 1.00 out of 1.00

What is a function of a modern data warehouse?

- a. supports batch processing only
- b. stores raw data only
- c. supports real-time and batch processing
- d. provides built-in or native online analytical processing

Your answer is correct.

The correct answer is: provides built-in or native online analytical processing

Question **139**

Correct

Mark 3.00 out of 3.00

DRAG DROP -

Your company plans to load data from a customer relationship management (CRM) system to a data warehouse by using an extract, load, and transform Where does data processing occur for each stage of the ELT process? To answer, drag the appropriate locations to the correct stages. Each location may be used more than once, or not at all.

Extract: The CRM system

Load: The data warehouse

Transform: The data warehouse

An in-memory data integration tool

The data warehouse

Your answer is correct.

The correct answer is:
DRAG DROP -

Your company plans to load data from a customer relationship management (CRM) system to a data warehouse by using an extract, load, and transform Where does data processing occur for each stage of the ELT process? To answer, drag the appropriate locations to the correct stages. Each location may be used more than once, or not at all.

Extract:[The CRM system]

Load:[The data warehouse]

Transform:[The data warehouse]

Question **140**

Correct

Mark 1.00 out of 1.00

A bank has a system that manages financial transactions.
When transferring money between accounts. the system must never retrieve a value for the source account that reflects the balance before the transfer and a value for the destination account that reflects the balance after the transfer.
Of which ACID semantic is this an example?

- a. isolation ✓
- b. consistency
- c. durability
- d. atomicity

Your answer is correct.

The correct answer is: isolation

Question **141**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

A view ✓ presents content defined by a query.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.
[A view ] presents content defined by a query.

Question **142**

Incorrect

Mark 0.00 out of 1.00

You have the following T-SQL statement.

SELECT o.OrderNo, o.OrderDate, c.Address, c.City

FROM Order AS o -

JOIN Customer AS c -
ON o.Customer = c.ID;

What is the function of the ON clause?

- a. to update data in the Customer table
- b. to define an alias for the Customer table
- c. to update data in the Order table
- d. to match data in the Customer table to data in the Order table

Your answer is incorrect.

The correct answer is: to match data in the Customer table to data in the Order table

Question **143**

Partially correct

Mark 2.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Copying data to Azure Data Lake Storage from another Azure region results in extra cost. | NO |
| You can use blob, table, and file storage in the same Azure Storage account. | YES |
| You implement Azure Data Lake Storage by creating an Azure Storage account. | YES |

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: Copying data to Azure Data Lake Storage from another Azure region results in extra cost. → YES, You can use blob, table, and file storage in the same Azure Storage account. → YES, You implement Azure Data Lake Storage by creating an Azure Storage account. → YES

Question **144**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Apache Spark Structured Streaming is built on a [dataframe] that is used to encapsulate a single table of data.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

Apache Spark Structured Streaming is built on a [dataframe] that is used to encapsulate a single table of data.

Question **145**

Correct

Mark 1.00 out of 1.00

What should you use to build a Microsoft Power BI paginated report?

a. Charticulator

b. Power BI Report Builder

c. the Power BI service

d. Power BI Desktop

Your answer is correct.

The correct answer is:
Power BI Report Builder

Question **146**

Correct

Mark 1.00 out of 1.00

In a fully denormalized database, how is data read and written for a single entity?

a. Data is read from multiple tables and written to multiple tables.

b. Data is read from multiple tables and written to a single table.

c. Data is read from a single table and written to multiple tables.

d. Data is read from a single table and written to a single table.

Your answer is correct.

The correct answers are: Data is read from a single table and written to a single table., Data is read from a single table and written to multiple tables.

Question **147**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

| | |
|---|---|
| In non-relational databases, data is always stored in tables. | NO |
| In non-relational databases, data is always stored by using a fixed schema. | NO |
| In non-relational databases, two entities in the same collection can have different sets of fields. | YES |

Your answer is correct.

The correct answer is: In non-relational databases, data is always stored in tables. → NO, In non-relational databases, data is always stored by using a fixed schema. → NO, In non-relational databases, two entities in the same collection can have different sets of fields. → YES

Question **148**

Incorrect

Mark 0.00 out of 1.00

Select the answer that correctly completes the sentence.

Microsoft SQL Server database tables store semi structured data .

Your answer is incorrect.

The correct answer is:
Select the answer that correctly completes the sentence.

Microsoft SQL Server database tables store [relational data].

Question **149**

Correct

Mark 1.00 out of 1.00

You have the following SQL query.

```
INSERT INTO dbo.Products (ProductID, ProductName, Price, ProductDescription)
    VALUES (1, 'Clamp', 12.48, 'Workbench clamp') ;
```

What are dbo.Products and ProductName? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Dbo. Products: A table

ProductName: A column

Your answer is correct.

The correct answer is:
You have the following SQL query.

```
INSERT INTO dbo.Products (ProductID, ProductName, Price, ProductDescription)
    VALUES (1, 'Clamp', 12.48, 'Workbench clamp') ;
```

What are dbo.Products and ProductName? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.
Dbo. Products:[A table ]

ProductName:[A column ]

Question **150**

Correct

Mark 1.00 out of 1.00

Your company needs to design a database that illustrates the relationships between utilization levels of individual network devices across a local area network.
Which type of data store should you use?

- a. columnar
- b. graph
- c. key/value
- d. document

Your answer is correct.

The correct answer is: graph

Question **151**

Incorrect

Mark 0.00 out of 1.00

SELECT, INSERT, and UPDATE are examples of which type of SQL statement?

- a. Data Control Language (DCL)
- b. Data Manipulation Language (DML)
- c. Data Definition Language (DDL) ✗

Your answer is incorrect.

The correct answer is: Data Manipulation Language (DML)

Question **152**

Correct

Mark 2.00 out of 2.00

In Azure Table storage, each row in a table must be uniquely identified by which two components? Each correct answer presents part of the solution.
NOTE: Each correct selection is worth one point

- [ ] a. a range
- [x] b. a partition key ✓
- [x] c. a row key ✓
- [ ] d. a timestamp

Your answer is correct.

The correct answers are: a partition key, a row key

Question **153**

Correct

Mark 1.00 out of 1.00

Your company is designing a data store that will contain student data. The data has the following format.

| StudentNumber | StudentInformation |
|---|---|
| 7634634 | First name: Ben<br>Last: Smith<br>Preferred Name: Benjamin |
| 7634634 | First Name: Dominik<br>Last Name: Paiha<br>Email Address: dpaiha@contoso.com<br>MCP ID: 931817 |
| 7634636 | First Name: Reshma<br>Last Name: Patel<br>Phone number: 514-555-1101 |
| 7634637 | First Name: Yun-Feng<br>Last Name: Peng |

Which type of data store should you use?

a. key/value

b. columnar

c. object

d. graph

Your answer is correct.

The correct answer is:

columnar

Question **154**

Incorrect

Mark 0.00 out of 1.00

Which Azure SQL offering provides near-100 percent compatibility with on-premises Microsoft SQL Server instances, while providing automated updates. backups, and maintenance tasks?

a. Azure SQL Managed Instance

b. Azure SQL Database

c. Azure SQL Edge

d. SQL Server on Azure Virtual Machines

Your answer is incorrect.

The correct answer is: Azure SQL Managed Instance

Question **155**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Data stored in the Hot access tier of Azure Blob storage is stored on high-performance media. Yes

The Cool access tier in Azure Blob storage incurs higher storage costs compared to the Hot access tier. No

The Archive access tier in Azure Blob storage has the highest latency of the three available tiers. Yes

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Data stored in the Hot access tier of Azure Blob storage is stored on high-performance media. [Yes]
The Cool access tier in Azure Blob storage incurs higher storage costs compared to the Hot access tier. [No]
The Archive access tier in Azure Blob storage has the highest latency of the three available tiers. [Yes]

Question **156**

Correct

Mark 1.00 out of 1.00

Which command-line tool can you use to query Azure SQL databases?

a. sqlcmd

b. bcp

c. Azure CLI

d. azdata

Your answer is correct.

The correct answer is: sqlcmd

Question **157**

Incorrect

Mark 0.00 out of 1.00

Which T-SQL statement should be used to instruct a database management system to use an index instead of performing a full table scan?

- a. WHERE
- b. SELECT
- c. JOIN ⊗

Your answer is incorrect.

**Correct Answer:** SELECT

**Explanation:**

- The **SELECT** statement is the one that **initiates a query** and is where index hints (such as WITH (INDEX(...))) are applied **if you want to explicitly instruct the DBMS to use a specific index**.
- Neither WHERE nor JOIN can **by themselves** force index usage — they can only **influence** the query optimizer's decision.
- **Forcing index use is only possible within the SELECT statement**, using **query hints**.

The correct answer is:
SELECT

Question **158**

Incorrect

Mark 0.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

An extract, transform, and load (ETL) process requires [a target data store powerful enough to transform data.] ⊗

Your answer is incorrect.

## Correct Answer:

**data that is fully processed before being loaded to the target data store.**

---

## Reason:

In an **ETL (Extract, Transform, Load)** process:

- **Data is first extracted** from various sources.
- It is **transformed** (cleaned, standardized, aggregated, etc.).
- Then the **fully processed data is loaded** into the target system.

This differs from **ELT**, where data is loaded first and transformed later.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

An extract, transform, and load (ETL) process requires [data that is fully processed before being loaded to the target data store.]

Question **159**

Correct

Mark 1.00 out of 1.00

Your company needs to ensure that an Azure virtual machine can connect to Azure SQL databases without exposing the databases to the internet. What should you use?

- a. Azure Traffic Manager
- b. Azure Application Gateway
- c. Azure DNS
- d. Azure Private Link

Your answer is correct.

The correct answer is:
Azure Private Link

Question **160**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Azure Databricks is the Azure implementation of a comprehensive data analytics solution that is built on Apache Spark and is available on multiple cloud platforms.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

[Azure Databricks] is the Azure implementation of a comprehensive data analytics solution that is built on Apache Spark and is available on multiple cloud platforms.

Question **161**

Incorrect

Mark 0.00 out of 1.00

Select the answer that correctly completes the sentence.

Key-value stores ✗ natively support the analysis of relationships between entities

Your answer is incorrect.

**Correct Answer:** Graph databases

**Reason:**

**Graph databases** are specifically designed to store and analyze **relationships** between entities using **nodes** (entities) and **edges** (relationships). They are ideal for use cases like:

- Social networks
- Fraud detection
- Organizational hierarchies
- Recommendation engines

They use languages like **Gremlin** or **Cypher** for querying relationships.

The correct answer is:
Select the answer that correctly completes the sentence.

[Graph databases ] natively support the analysis of relationships between entities

Question **162**

Correct

Mark 1.00 out of 1.00

Which statement is an example of Data Definition Language (DDL)?

a. SELECT

b. JOIN

c. CREATE ✓

d. UPDATE

Your answer is correct.

The correct answer is:
CREATE

Question **163**

Correct

Mark 1.00 out of 1.00

A treemap is a chart of colored rectangles. The size of the rectangles represent the relative value of each item. The chart can be hierarchical, displaying data as a set of rectangles nested within the main rectangle.

Your answer is correct.

The correct answer is:
A [treemap] is a chart of colored rectangles. The size of the rectangles represent the relative value of each item. The chart can be hierarchical, displaying data as a set of rectangles nested within the main rectangle.

Question **164**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

A block of code that runs in a database is called

- a. a view.
- b. an index.
- c. a stored procedure.
- d. a table.

Your answer is correct.

The correct answer is: a stored procedure.

Question **165**

Correct

Mark 1.00 out of 1.00

You need to create a visualization of running sales totals per quarter as shown in the following exhibit.

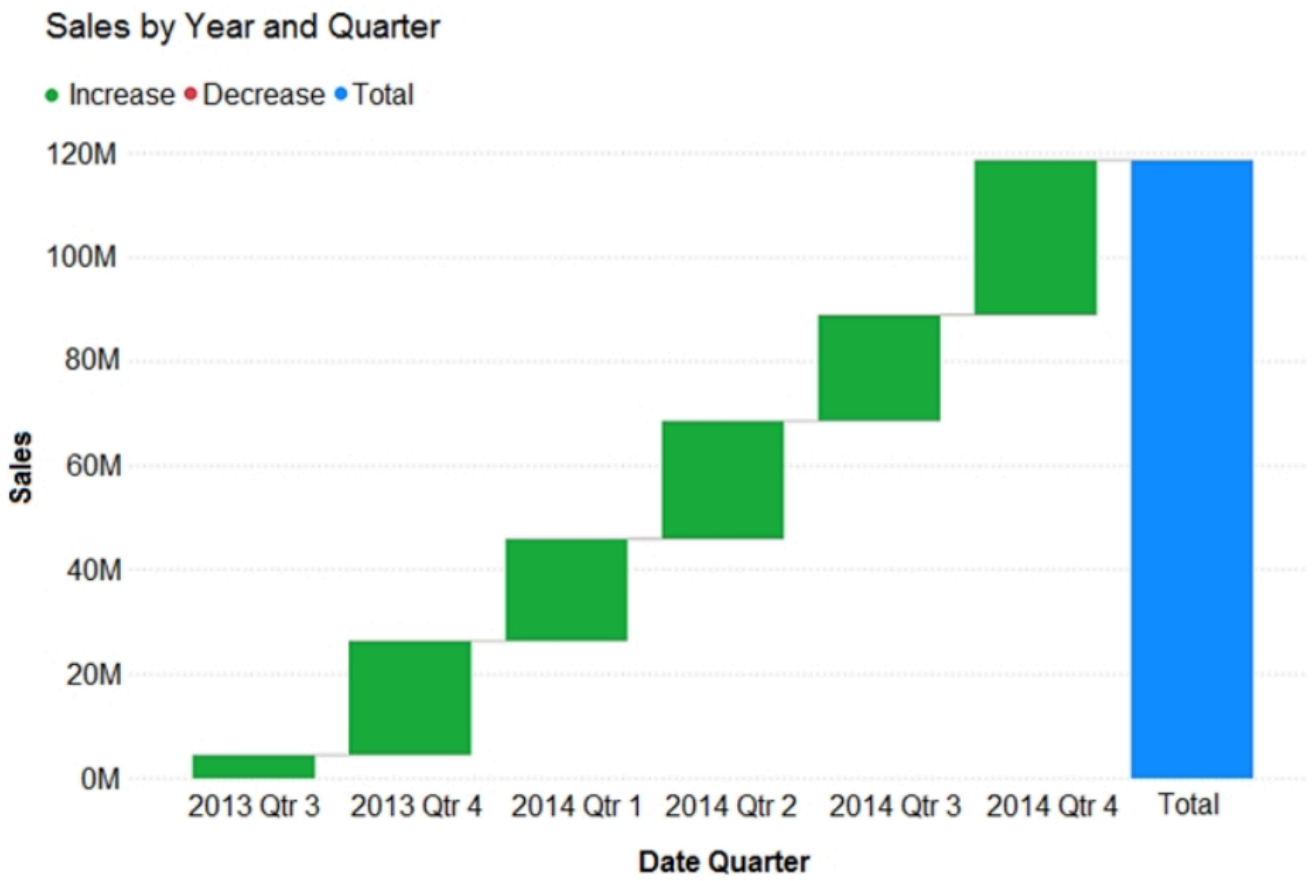


What should you create in Power BI Desktop?

- a. a waterfall chart
- b. a ribbon chart
- c. a bar chart
- d. a decomposition tree

Your answer is correct.

The correct answer is: a waterfall chart

Question **166**

Correct

Mark 1.00 out of 1.00

In an analytical data model, which type of table contains entities that are used to aggregate numeric values, where each entity is represented by a row with a unique key value?

- a. bridge
- b. fact
- c. dimension

Your answer is correct.

The correct answer is: fact

Question **167**

Correct

Mark 3.00 out of 3.00

Match the job roles to the appropriate tasks.

To answer, drag the appropriate job role from the column on the left to its task on the right. Each role may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

| | |
|---|---|
| Data engineer | Create a pipeline to transfer data from a database to a data warehouse. |
| Database administrator | Create a transactional database. |
| Data analyst | Create a chart that shows sales transactions by week. |

Data architect

Your answer is correct.

The correct answer is:
Match the job roles to the appropriate tasks.

To answer, drag the appropriate job role from the column on the left to its task on the right. Each role may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

[Data engineer] Create a pipeline to transfer data from a database to a data warehouse.
[Database administrator] Create a transactional database.
[Data analyst] Create a chart that shows sales transactions by week.

Question **168**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

You create Azure File Storage in a storage account. Yes

The Premium tier of Azure File Storage uses HDD-based hardware in a datacenter. No

The Standard tier of Azure File Storage uses SSDs. No

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

You create Azure File Storage in a storage account. [Yes]
The Premium tier of Azure File Storage uses HDD-based hardware in a datacenter. [No]
The Standard tier of Azure File Storage uses SSDs. [No]

Question **169**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

A relational database is used to serve structured data in tables.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

A [relational] database is used to serve structured data in tables.

Question **170**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

In Azure Blob Storage, each block in a block blob is a fixed size. No

In an Azure Blob Storage account, blobs are stored in a container. Yes

In an Azure Blob Storage account that does NOT have a hierarchical namespace enabled, you can assign access at the folder level. No

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

In Azure Blob Storage, each block in a block blob is a fixed size. [No]
In an Azure Blob Storage account, blobs are stored in a container. [Yes]
In an Azure Blob Storage account that does NOT have a hierarchical namespace enabled, you can assign access at the folder level. [No]

Question **171**

Incorrect

Mark 0.00 out of 1.00

Which activity is most common for transactional workloads?

a. aggregating massive amounts of data

b. producing complex reports

c. performing self-service data exploration

d. recording small units of work events in real time

Your answer is incorrect.

The correct answer is: recording small units of work events in real time

Question **172**

Correct

Mark 3.00 out of 3.00

DRAG DROP -
Match the terms to the appropriate descriptions.
To answer, drag the appropriate term from the column. Each term may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Table ✓ A database object that holds data

View ✓ A database object whose content is defined by a query

Index ✓ A database object that helps improve the speed of data retrieval

Select and Place:

Your answer is correct.

The correct answer is:
DRAG DROP -
Match the terms to the appropriate descriptions.
To answer, drag the appropriate term from the column. Each term may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

[Table]A database object that holds data

[View]A database object whose content is defined by a query

[Index]A database object that helps improve the speed of data retrieval

Select and Place:

Question **173**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

An elastic pool ✓ enables multiple Azure SQL databases to share resources such as memory, data storage space, and processing power.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

[An elastic pool] enables multiple Azure SQL databases to share resources such as memory, data storage space, and processing power.

Question **174**

Incorrect

Mark 0.00 out of 1.00

You have structured data in tabular format.

What represents an individual instance of a data entity?

- a. a column
- b. a table
- c. a cell
- d. a row

Your answer is incorrect.

The correct answer is: a row

Question **175**

Incorrect

Mark 0.00 out of 1.00

Select the answer that correctly completes the sentence.

Azure SQL Database is a platform as a service (PaaS) relational data service that uses massively parallel processing (MPP).

Your answer is incorrect.

**Correct Answer:** Azure Synapse Analytics

## Explanation:

**Azure Synapse Analytics** is a **Platform as a Service (PaaS)** offering that supports **massively parallel processing (MPP)**, making it ideal for:

- Running complex queries across large volumes of data
- Performing big data analytics
- Serving as a modern cloud data warehouse

The correct answer is:
Select the answer that correctly completes the sentence.

[Azure Synapse Analytics ] is a platform as a service (PaaS) relational data service that uses massively parallel processing (MPP).

Question **176**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Data engineers identify data governance rules. No

Data engineers create data ingestion pipelines. Yes

Data engineers implement backup and recovery processes for database systems. No

Your answer is correct.

The correct answer is:

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Data engineers identify data governance rules. [No]
Data engineers create data ingestion pipelines. [Yes]
Data engineers implement backup and recovery processes for database systems. [No]

Question **177**

Correct

Mark 1.00 out of 1.00

What is the typical latency for read operations against media files stored in the Cool access tier of Azure Blob Storage?

a. a few milliseconds

b. a few hours

c. a few seconds

d. a few minutes

Your answer is correct.

The correct answer is: a few milliseconds

Question **178**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completed the sentence.

In Microsoft Power BI, you can create [an attribute hierarchy] to quickly drill-up or drill-down to find aggregated values at different levels of a dimension.

Your answer is correct.

The correct answer is:
Select the answer that correctly completed the sentence.

In Microsoft Power BI, you can create [an attribute hierarchy] to quickly drill-up or drill-down to find aggregated values at different levels of a dimension.

Question **179**

Incorrect

Mark 0.00 out of 1.00

Blob rehydration occurs in Azure Blob storage when a blob moves between which access tiers?

a. Archive to Cool

b. Hot to Archive

c. Hot to Cool

d. Cool to Archive

Your answer is incorrect.

The correct answer is: Archive to Cool

Question **180**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

The Azure Cosmos Core (SQL) API enables the use of SELECT statements toretrieve documents from Azure Cosmos DB.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

The Azure Cosmos [Core (SQL) ] API enables the use of SELECT statements toretrieve documents from Azure Cosmos DB.

Question **181**

Partially correct

Mark 2.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| You can copy a dashboard between Microsoft Power BI workspaces. | YES |
| A Microsoft Power BI dashboard can only display visualizations from a single dataset. | NO |
| A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook. | YES |

Your answer is partially correct.

You have correctly selected 2.

The correct answer is: You can copy a dashboard between Microsoft Power BI workspaces.

→ NO, A Microsoft Power BI dashboard can only display visualizations from a single dataset. → NO,
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook. → YES

Question **182**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

You can query a graph database in Azure Cosmos DB as nodes and edges by using the Gremlin language

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

You can query a graph database in Azure Cosmos DB [as nodes and edges by using the Gremlin language]

Question **183**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

The pipeline engine used in Azure Data Factory is also available in Azure Synapse Analytics.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

The pipeline engine used in Azure Data Factory is also available in [Azure Synapse Analytics.]

Question **184**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| You must apply operating system updates to Azure SQL databases regularly. | NO |
| You need a Microsoft 365 subscription to create an Azure SQL database. | NO |
| You can use existing Microsoft SQL Server licenses to reduce the cost of Azure SQL databases. | YES |

Your answer is correct.

The correct answer is: You must apply operating system updates to Azure SQL databases regularly. → NO, You need a Microsoft 365 subscription to create an Azure SQL database. → NO, You can use existing Microsoft SQL Server licenses to reduce the cost of Azure SQL databases. → YES

Question **185**

Correct

Mark 1.00 out of 1.00

In a fully normalized database, how is data read and written for a single entity?

- a. Data is read from a single table and written to multiple tables.
- b. Data is read from multiple tables and written to a single table. ✓
- c. Data is read from a single table and written to a single table.
- d. Data is read from multiple tables and written to multiple tables.

Your answer is correct.

The correct answer is: Data is read from multiple tables and written to a single table.

Question **186**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

A visualization that illustrates a university's current student enrolment per department is an example of [descriptive] ✓ analytics.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.
A visualization that illustrates a university's current student enrolment per department is an example of [descriptive ] analytics.

Question **187**

Correct

Mark 1.00 out of 1.00

Which statement is an example of Data Definition Language (DDL)?

- a. SELECT
- b. INSERT
- c. DROP ✓
- d. DELETE

Your answer is correct.

The correct answer is:
DROP

Question **188**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Creating closed caption text for audio files is an example of cognitive analytics.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.
Hot Area:

Creating closed caption text for audio files is an example of [cognitive] analytics.

Question **189**

Correct

Mark 3.00 out of 3.00

Match the job roles to the appropriate tasks.

To answer, drag the appropriate role from the column on the left to its task on the right. Each role may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

| | |
|---|---|
| Database administrator | Assign permissions to users. |
| Data analyst | Derive meaning from data by creating visualizations and charts. |
| Data Engineer | Implement which pipelines to transfer and transform data between systems. |

Your answer is correct.

The correct answer is:
Match the job roles to the appropriate tasks.

To answer, drag the appropriate role from the column on the left to its task on the right. Each role may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

[Database administrator] Assign permissions to users.
[Data analyst] Derive meaning from data by creating visualizations and charts.
[Data Engineer] Implement which pipelines to transfer and transform data between systems.

Question **190**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

A clustered index is an object associated with a table that sorts on their key values.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

[A clustered index] is an object associated with a table that sorts on their key values.

Question **191**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

Batch workloads collect data and then process the data when a condition is met.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

Batch workloads [collect data and then process the data when a condition is met.]

Question **192**

Correct

Mark 1.00 out of 1.00

Which Azure Cosmos DB API should you use for a graph database?

- a. Gremlin
- b. Table
- c. NoSQL
- d. Cassandra

Your answer is correct.

The correct answer is: Gremlin

Question **193**

Correct

Mark 1.00 out of 1.00

- Select the answer that correctly completes the sentence.

In a SQL statement, MAX and MIN are examples of aggregate functions.

Your answer is correct.

The correct answer is:

- Select the answer that correctly completes the sentence.

In a SQL statement,[ MAX and MIN] are examples of aggregate functions.

Question **194**

Partially correct

Mark 2.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

The pgAdmin tool can be used to manage Azure Database for PostgreSQL. Yes

Azure Database for MariaDB offers point-in-time restore for up to 365 days. No

Azure Database for MySQL always requires SSL connections. Yes

Your answer is partially correct.

You have correctly selected 2.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

The pgAdmin tool can be used to manage Azure Database for PostgreSQL. [Yes]
Azure Database for MariaDB offers point-in-time restore for up to 365 days. [No]
Azure Database for MySQL always requires SSL connections. [No]

Question **195**

Correct

Mark 1.00 out of 1.00

Which Azure storage solution provides native support for POSIX-compliant access control lists (ACLs)?

a. Azure Data Lake Storage

b. Azure Files

c. Azure Table storage

d. Azure Queue storage

Your answer is correct.

The correct answer is:
Azure Data Lake Storage

Question **196**

Incorrect

Mark 0.00 out of 1.00

Your company needs to implement a relational database in Azure. The solution must minimize ongoing maintenance.
Which Azure service should you use?

- a. SQL Server on Azure Virtual Machines
- b. Azure Cosmos DB
- c. Azure HDInsight
- d. Azure SQL Database

Your answer is incorrect.

**Correct Answer:**Azure SQL Database

**Reason:**

**Azure SQL Database** is a **fully managed PaaS** service that:

- Minimizes ongoing maintenance by handling **patching, backups, scaling, and high availability**,
- Supports **relational workloads** with T-SQL and built-in intelligence,
- Is ideal when you want to avoid managing infrastructure or SQL Server instances yourself.

The correct answer is:
Azure SQL Database

Question **197**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a graph.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a [graph. ]

Question **198**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

In an analytical model, any point where the dimensions intersect represents an aggregated measure.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

In an analytical model, any point where the dimensions intersect represents [an aggregated measure.]

Question **199**

Correct

Mark 3.00 out of 3.00

Match the terms to the appropriate descriptions.

To answer, drag the appropriate term from the column on the left to its description on the right. Each term may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

| | |
|---|---|
| Extract, transform, and load (ETL) | A solution that copies data from files and online transaction processing (OLTP) databases into an analytical processing system. |
| Fact tables | Contains numeric values that can be analyzed. |
| Data warehouses | A solution that stores data in a relational schema that is optimized for read operations. |

Dimension tables

Online analytical processing (OLAP)

Your answer is correct.

The correct answer is:
Match the terms to the appropriate descriptions.

To answer, drag the appropriate term from the column on the left to its description on the right. Each term may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

[Extract, transform, and load (ETL)] A solution that copies data from files and online transaction processing (OLTP) databases into an analytical processing system.
[Fact tables] Contains numeric values that can be analyzed.
[Data warehouses] A solution that stores data in a relational schema that is optimized for read operations.

Question **200**

Incorrect

Mark 0.00 out of 1.00

Your company is designing an application that will write a high volume of JSON data and will have an application-defined schema. Which type of data store should you use?

- a. columnar
- b. key/value ⊗
- c. document
- d. graph

Your answer is incorrect.

**Correct Answer: Document**

**Explanation:**

You should use a **Document data store** because:

- It is **ideal for storing JSON data**.
- It supports **flexible, application-defined schemas** — meaning each document (like a JSON object) can have different fields and structures.
- It is optimized for **high-volume writes and reads** of semi-structured data.
- It allows **querying the data efficiently** using indexes on document fields.

The correct answer is: document

Question **201**

Correct

Mark 3.00 out of 3.00

Match the types of activities to the appropriate Azure Data Factory activities.
To answer, drag the appropriate activity type from the column to its Data Factory activity . Each activity type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

| | |
|---|---|
| Data movement | Copy |
| Data transformation | Mapping data flow |
| Control | Lookup |

Select and Place:

Your answer is correct.

The correct answer is:
Match the types of activities to the appropriate Azure Data Factory activities.
To answer, drag the appropriate activity type from the column to its Data Factory activity . Each activity type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

[Data movement]Copy
[Data transformation]Mapping data flow
[Control]Lookup

Select and Place:

Question **202**

Incorrect

Mark 0.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

A relational database is appropriate for scenarios that involve a high volume of changes to relationships between entities

Your answer is incorrect.

## Correct Answer:

**transactional writes**

## Reason:

Relational databases are optimized for **transactional workloads**, which involve frequent **inserts, updates, and deletes**—typically using **structured schemas and ACID compliance** to ensure data integrity. This makes them ideal for scenarios like order processing, financial systems, and customer records management.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

A relational database is appropriate for scenarios that involve a high volume of [transactional writes]

Question **203**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

A AVRO is a row-based file format where each record contains a JSON header and binary data.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

A [AVRO] is a row-based file format where each record contains a JSON header and binary data.

Question **204**

Correct

Mark 1.00 out of 1.00

What is used to define a query in a stream processing jobs in Azure Stream Analytics?

a. YAML

b. XML

c. KQL

d. SQL

Your answer is correct.

The correct answer is: SQL

Question **205**

Correct

Mark 1.00 out of 1.00

GRANT, REVOKE, and DENY are examples of which type of SQL statement?

a. Data Manipulation Language (DML)

b. Data Definition language (DDL)

c. Data Control Language (DCL)

Your answer is correct.

The correct answer is: Data Control Language (DCL)

Question **206**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Relational database tables contain columns and rows | YES |
| Indexes in a relational database describe the data types in a table | NO |
| A database view is a virtual table whose content is defined by a query | YES |

Your answer is correct.

The correct answer is: Relational database tables contain columns and rows → YES, Indexes in a relational database describe the data types in a table → NO, A database view is a virtual table whose content is defined by a query → YES

Question **207**

Correct

Mark 1.00 out of 1.00

You have the following two tables of customer data.

| CustomerID | NameStyle | Title | FirstName | MiddleName | LastName |
|---|---|---|---|---|---|
| 29486 | 0 | Ms. | Kim | NULL | Abercrombie |
| 29489 | 0 | Ms. | Frances | B. | Adams |
| 29490 | 0 | Ms. | Margaret | J. | Smith |
| 29492 | 0 | Mr. | Jay | NULL | Adams |

| CustomerID | AddressID | AddressType | ModifiedDate |
|---|---|---|---|
| 29486 | 621 | Main Office | 9/1/05 0:00 |
| 29489 | 1069 | Main Office | 7/1/05 0:00 |
| 29490 | 887 | Main Office | 9/1/06 0:00 |
| 29492 | 618 | Main Office | 12/1/06 0:00 |
| 29494 | 537 | Main Office | 9/1/05 0:00 |

Which type of data do the tables contain?

a. structured

b. non-relational

c. unstructured

d. semi-structured

Your answer is correct.

The correct answer is: structured

Question **208**

Correct

Mark 2.00 out of 2.00

You plan to deploy a PostgreSQL database to Azure.
Which hosting model corresponds to the available deployment options? To answer, select the appropriate options in the answer area.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| PostgreSQL on Azure VMs: | Infrastructure as a service (IaaS) |
| Azure Database for PostgreSQL: | Platform as a service (PaaS) |

Your answer is correct.

The correct answer is:
PostgreSQL on Azure VMs: → Infrastructure as a service (IaaS),

Azure Database for PostgreSQL: → Platform as a service (PaaS)

Question **209**

Correct

Mark 1.00 out of 1.00

A JSON document that contains customer entities with a variable number of properties is an example of semi-structured data.

Your answer is correct.

The correct answer is:
A JSON document that contains customer entities with a variable number of properties is an example of [semi-structured] data.

Question **210**

Correct

Mark 3.00 out of 3.00

DRAG DROP -
Your company plans to load data from a customer relationship management (CRM) system to a data warehouse by using an extract, load, and transform (ELT) process.
Where does data processing occur for each stage of the ELT process? To answer, drag the appropriate locations to the correct stages. Each location may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.
Select and Place:

Extract: The CRM system

Load: The data warehouse

Transform: The data warehouse

The data warehouse | A standalone data analysis tool | The CRM system

Your answer is correct.

The correct answer is:
DRAG DROP -
Your company plans to load data from a customer relationship management (CRM) system to a data warehouse by using an extract, load, and transform (ELT) process.
Where does data processing occur for each stage of the ELT process? To answer, drag the appropriate locations to the correct stages. Each location may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.
Select and Place:

Extract: [The CRM system]

Load:[The data warehouse]

Transform:[The data warehouse]

Question **211**

Correct

Mark 1.00 out of 1.00

You have data in a tabular format as shown in the following table.

| ID | First name | Last name | Email |
|---|---|---|---|
| 1000 | Ben | Smith | bensmith@contoso.com |
| 1001 | Carlos | Grilo | carlosgrilo@contoso.com |

Of which type of data is this an example?

a. structured data
b. document
c. unstructured data
d. semi-structured data

Your answer is correct.

The correct answer is: structured data

Question **212**

Correct

Mark 1.00 out of 1.00

Which Azure Data Factory component initiates the execution of a pipeline?

- a. an activity
- b. a control flow
- c. a trigger ✓
- d. a parameter

Your answer is correct.

The correct answer is: a trigger

Question **213**

Incorrect

Mark 0.00 out of 1.00

Which Azure Data Factory component provides the compute environment for activities?

- a. SSIS packages
- b. an integration runtime
- c. a pipeline
- d. a control flow ✗

Your answer is incorrect.

The correct answer is: **an integration runtime**

**Explanation:**

| Component | Purpose |
|---|---|
| **Integration Runtime (IR)** | Provides the **compute environment** to execute activities in Azure Data Factory, such as data movement, data flow, and external SSIS execution. |

The correct answer is:
an integration runtime

Question **214**

Correct

Mark 1.00 out of 1.00

You manage an application that stores data in a shared folder on a Windows server.
You need to move the shared folder to Azure Storage.
Which type of Azure Storage should you use?

- a. queue
- b. blob
- c. table
- d. file

Your answer is correct.

The correct answer is: file

Question **215**

Correct

Mark 1.00 out of 1.00

What can you use to create data ingestion pipelines?

- a. Azure HDInsight
- b. Azure Data Factory
- c. Azure SQL Database
- d. Azure Cosmos DB

Your answer is correct.

The correct answer is: Azure Data Factory

Question **216**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

Descriptive analytics tells you what occurred in the past.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

Descriptive analytics tells you [what occurred in the past.]

Question **217**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

The Cool access tier is the default tier of an Azure Storage account. No

Append blobs are optimized for random read and write operations. No

Retrieving data from the Archive access tier takes hours to complete.. Yes

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

The Cool access tier is the default tier of an Azure Storage account. [No]
Append blobs are optimized for random read and write operations. [No]
Retrieving data from the Archive access tier takes hours to complete.. [Yes]

Question **218**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

PolyBase is an Azure Synapse Analytics feature that retrieves data from Azure Blob storage by using Transact-SQL queries.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

[PolyBase] is an Azure Synapse Analytics feature that retrieves data from Azure Blob storage by using Transact-SQL queries.

Question **219**

Correct

Mark 1.00 out of 1.00

You are deploying a software as a service (SaaS) application that requires a relational database for Online Transaction Processing (OLTP). Which Azure service should you use to support the application?

- a. Azure SQL Database
- b. Azure Cosmos DB
- c. Azure HDInsight
- d. Azure Synapse Analytics

Your answer is correct.

The correct answer is:
Azure SQL Database

Question **220**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

A clustered index physically sorts the data in a table based on the values in a specified column.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.
[A clustered index ] physically sorts the data in a table based on the values in a specified column.

Question **221**

Incorrect

Mark 0.00 out of 1.00

Your company recently reported sales from the third quarter.
You have the chart shown in the following exhibit.

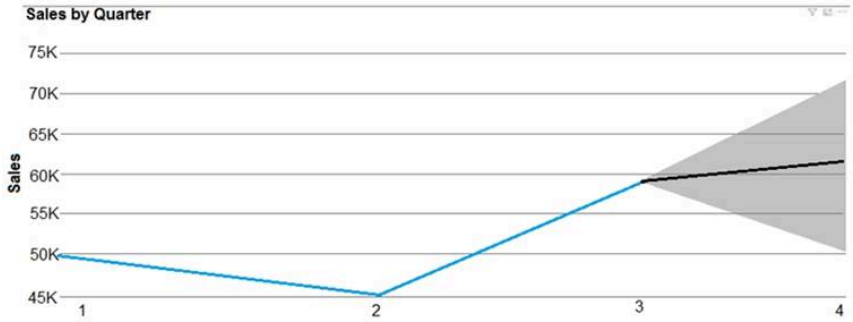


Which type of analysis is shown in the fourth quarter?

- a. prescriptive ✗
- b. predictive
- c. diagnostic
- d. descriptive

Your answer is incorrect.

**Correct Answer:** Predictive

**Reason:**

The chart shows actual sales data for Q1 to Q3, and for Q4, it includes:

- A **forecasted trend line** (black line),
- A **shaded area** representing a **range of possible values** (confidence interval).

This is a classic example of **predictive analytics**, which uses historical data to **predict future outcomes**.

The correct answer is: predictive

Question **222**

Correct

Mark 1.00 out of 1.00

Which statement is an example of Data Manipulation Language (DML)?

- a. DISABLE
- b. GRANT
- c. INSERT ✓
- d. REVOKE

Your answer is correct.

The correct answer is: INSERT

Question **223**

Correct

Mark 1.00 out of 1.00

You need to store data in Azure Blob storage for seven years to meet your company's compliance requirements. The retrieval time of the data is unimportant. The solution must minimize storage costs.
Which storage tier should you use?

- a. Hot
- b. Cool
- c. Archive

Your answer is correct.

The correct answer is:
Archive

Question **224**

Incorrect

Mark 0.00 out of 1.00

A data engineer is responsible for which task?

- a. Explore data to identify trends.
- b. Design and implement data stores for analytical workloads.
- c. Design and build analytical models.
- d. Implement policies, tools, and processes for backup and recovery plans.

Your answer is incorrect.

The correct answer is: Design and implement data stores for analytical workloads.

Question **225**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

In a data warehousing workload, data [from multiple sources is combined in a single location]

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

In a data warehousing workload, data [from multiple sources is combined in a single location]

Question **226**

Correct

Mark 1.00 out of 1.00

When provisioning an Azure Cosmos DB account, which feature provides redundancy within an Azure region?

- a. multi-master replication
- b. the strong consistency level
- c. Availability Zones
- d. automatic failover

Your answer is correct.

The correct answer is:
Availability Zones

Question **227**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Azure Data Explorer uses KQL to explore, filter, and aggregate data quickly.

Your answer is correct.

The correct answer is:

Select the answer that correctly completes the sentence.

Azure Data Explorer uses [KQL] to explore, filter, and aggregate data quickly.

Question **228**

Correct

Mark 1.00 out of 1.00

You need to develop a solution to provide data to executives. The solution must provide an interactive graphical interface, depict various key performance indicators, and support data exploration by using drill down.
What should you use in Microsoft Power BI?

a. a report

b. a dataflow

c. a view

d. Microsoft Power Apps

Your answer is correct.

The correct answer is:
a report

Question **229**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

Transparent Data Encryption (TDE) encrypts [the database to protect data at rest.]

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

Transparent Data Encryption (TDE) encrypts [the database to protect data at rest.]

Question **230**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| Extract, transform, and load (ETL) can reduce the transfer of sensitive data to destination systems. | YES |
| Extract, load, and transform (ELT) transforms data by using a compute resource independent of the source system and destination system. | NO |
| Extract, load, and transform (ELT) minimizes the time it takes to copy large volumes of data to destination systems. | YES |

Your answer is correct.

The correct answer is: Extract, transform, and load (ETL) can reduce the transfer of sensitive data to destination systems. → YES, Extract, load, and transform (ELT) transforms data by using a compute resource independent of the source system and destination system. → NO, Extract, load, and transform (ELT) minimizes the time it takes to copy large volumes of data to destination systems. → YES

Question **231**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Azure Synapse Analytics combines a scalable, high-performance Microsoft SQL Server-based relational data warehouse with the flexibility of a data lake and open-source Apache Spark.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

[Azure Synapse Analytics] combines a scalable, high-performance Microsoft SQL Server-based relational data warehouse with the flexibility of a data lake and open-source Apache Spark.

Question **232**

Correct

Mark 1.00 out of 1.00

Which statement is an example of Data Definition Language (DDL)?

- a. INSERT
- b. MERGE
- c. CREATE ✓
- d. SELECT

Your answer is correct.

The correct answer is:
CREATE

Question **233**

Correct

Mark 2.00 out of 2.00

DRAG DROP -
You have a table named Sales that contains the following data.

| SalesDate | SalesAmount | ProductID |
|---|---|---|
| 4-Apr-20 | $2,000 | 1 |
| 5-Apr-20 | $40 | 2 |
| 5-Apr-20 | $2,300 | 1 |
| 6-Apr-20 | $40 | 3 |
| 6-Apr-20 | $200 | 4 |

You need to query the table to return the average sales amount per day. The output must produce the following results.

| SalesDate | AVG(SalesAmount) |
|---|---|
| 4-Apr-20 | $2,000 |
| 5-Apr-20 | $1,170 |
| 6-Apr-20 | $120 |

How should you complete the query? To answer, drag the appropriate values to the correct targets. Each value may be used once, more than once, or not at all.
You may need to drag the split bar between panes or scroll to view content.
NOTE: Each correct selection is worth one point.
Select and Place:

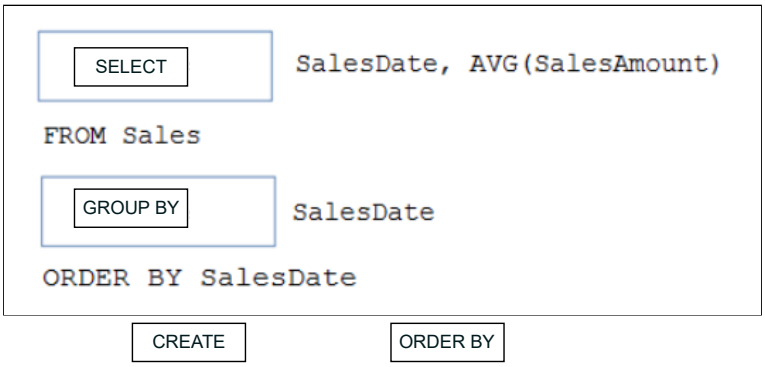


Your answer is correct.

Question **234**

Correct

Mark 1.00 out of 1.00

Which Azure Cosmos DB API should you use for a graph database?

- a. Azure Cosmos DB for Apache Cassandra
- b. Azure Cosmos DB for Table
- c. Azure Cosmos DB for Apache Gremlin
- d. Azure Cosmos DB for NoSQL

Your answer is correct.

The correct answer is: Azure Cosmos DB for Apache Gremlin

Question **235**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Online transaction processing (OLTP) relies on a database system where data storage is optimized for both read and write operations. Yes

Online analytical processing (OLAP) relies on a database system where data storage is optimized for both read and write operations. No

Data warehouses are optimized for read operations. Yes

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Online transaction processing (OLTP) relies on a database system where data storage is optimized for both read and write operations. [Yes]
Online analytical processing (OLAP) relies on a database system where data storage is optimized for both read and write operations. [No]
Data warehouses are optimized for read operations. [Yes]

Question **236**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

In an analytical model designed as a star schema, the entities by which you aggregate or slice measures are called dimensions.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

In an analytical model designed as a star schema, the entities by which you aggregate or slice measures are called [dimensions.]

Question **237**

Correct

Mark 1.00 out of 1.00

A team of developers has computers that run Windows 10 and Ubuntu Desktop.
The developers need to connect to and query an Azure SQL database from each of their computers. The developers require code assistance features such as
IntelliSense.
What should the developers use?

a. Azure Data Explorer

b. Azure Data Studio

c. sqlcmd

d. Microsoft SQL Server Management Studio (SSMS)

Your answer is correct.

The correct answer is:
Azure Data Studio

Question **238**

Correct

Mark 1.00 out of 1.00

What is a difference between structured data and semi-structured data?

- a. Structured data has a flexible schema and semi-structured data has a fixed schema.
- b. Structured data has a fixed schema and semi-structured data has a flexible schema.
- c. Only Structured data supports attributes.
- d. Only Structured data supports entities.

Your answer is correct.

The correct answer is: Structured data has a fixed schema and semi-structured data has a flexible schema.

Question **239**

Correct

Mark 1.00 out of 1.00

Which type of database is Azure Database for PostgreSQL?

- a. Microsoft SQL Server
- b. Platform as a service (PaaS)
- c. on-premises
- d. Infrastructure as a service (IaaS)

Your answer is correct.

The correct answer is:
Platform as a service (PaaS)

Question **240**

Incorrect

Mark 0.00 out of 1.00

You need to recommend a non-relational data store that is optimized for storing and retrieving text files, videos, audio streams, and virtual disk images. The data store must store data, some metadata, and a unique ID for each file.
Which type of data store should you recommend?

- a. columnar
- b. document
- c. key/value ⊗
- d. object

Your answer is incorrect.

**Correct Answer:** Object

**Explanation:**

You should recommend an **Object data store** because it is **specifically optimized** for:

- Storing **large binary files** such as:
  - 📁 Text files
  - 🎥 Videos
  - 🎧 Audio streams
  - 💾 Virtual disk images
- Storing data **along with metadata and a unique identifier** for each object (file)
- Supporting **HTTP-based access** and **global scalability** (as seen in Azure Blob Storage or Amazon S3)

The correct answer is: object

Question **241**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

An Azure Storage account must have Hierarchical namespace enabled to support Azure Data Lake Storage.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

An Azure Storage account must have [Hierarchical namespace] enabled to support Azure Data Lake Storage.

Question **242**

Correct

Mark 3.00 out of 3.00

DRAG DROP -
Match the types of data stores to the appropriate scenarios.
To answer, drag the appropriate data store type from the column on the left to its scenario on the right. Each data store type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Key/value ✓ A list of application user names and, for each, the default language

Object ✓ A collection of medical images and their associated metadata.

Graph ✓ Employee data that shows the relationships between employees

Select and Place:

Your answer is correct.

The correct answer is:
DRAG DROP -
Match the types of data stores to the appropriate scenarios.
To answer, drag the appropriate data store type from the column on the left to its scenario on the right. Each data store type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

[Key/value]A list of application user names and, for each, the default language

[Object]A collection of medical images and their associated metadata.

[Graph]Employee data that shows the relationships between employees

Select and Place:

Question **243**

Correct

Mark 1.00 out of 1.00

Your company has a reporting solution that has paginated reports. The reports query a dimensional model in a data warehouse.
Which type of processing does the reporting solution use?

a. batch processing

b. Online Transaction Processing (OLTP)

c. stream processing

d. Online Analytical Processing (OLAP) ✓

Your answer is correct.

The correct answer is:
Online Analytical Processing (OLAP)

Question **244**

Incorrect

Mark 0.00 out of 1.00

Which type of Azure resource supports the serverless configuration of an Azure SQL database?

- a. an Azure SQL Database elastic pool
- b. SQL Server on Azure Virtual Machines
- c. Azure SQL Managed Instance
- d. a single database in Azure SQL Database

Your answer is incorrect.

The correct answer is: a single database in Azure SQL Database

Question **245**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

a non-relational database [stores data without a fixed schema]

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

a non-relational database [stores data without a fixed schema]

Question **246**

Correct

Mark 3.00 out of 3.00

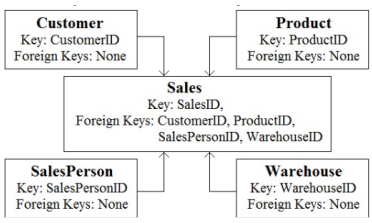


Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

The data model is a star schema .

Customer is a dimension table.

Your answer is correct.

The correct answer is:

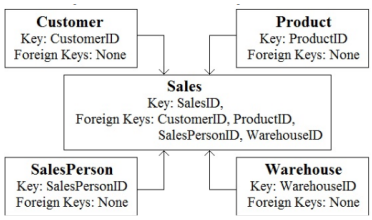


Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

The data model is a [star schema].

Customer is a [dimension]table.

Question **247**

Correct

Mark 1.00 out of 1.00

You need to use Transact-SQL to query files in Azure Data Lake Storage Gen 2 from an Azure Synapse Analytics data warehouse.
What should you use to query the files?

a. PolyBase

b. Azure Functions

c. Microsoft SQL Server Integration Services (SSIS)

d. Azure Data Factory

Your answer is correct.

The correct answer is:
PolyBase

Question **248**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

A key/value data store is optimized for simple lookups

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

A key/value data store is optimized for [simple lookups]

Question **249**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

| | |
|---|---|
| Batch processing can output data to a file store | YES |
| Batch processing can output data to a relational database | YES |
| Batch processing can output data to a NoSQL database | YES |

Your answer is correct.

The correct answer is: Batch processing can output data to a file store → YES, Batch processing can output data to a relational database → YES, Batch processing can output data to a NoSQL database → YES

Question **250**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

In batch processing, latency in delivering data processing results is acceptable.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

In batch processing,[latency in delivering data processing results is acceptable.]

Question **251**

Correct

Mark 1.00 out of 1.00

You have data saved in the following format.

```
FirstName, LastName, Age, LeisureHobby, SportsHobby
John, Smith, 23, Reading, Backteball
Ben, Smith, 21, Guitar Curling
```

Which format was used?

- a. JSON
- b. YAML
- c. HTML
- d. CSV ✓

Your answer is correct.

The correct answer is: CSV

Question **252**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Varying fields for each entity in a JSON document is an example of semi-structured data. ✓ .

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.
Varying fields for each entity in a JSON document is an example of [semi-structured data. ].

Question **253**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

A JSON document is an example of semi-structured data. ✓

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

A JSON document is an example of [semi-structured data.]

Question **254**

Correct

Mark 1.00 out of 1.00

Which setting can only be configured during the creation of an Azure Cosmos DB account?

- a. API
- b. geo-redundancy
- c. multi-region writes
- d. production or non-production account type

Your answer is correct.

The correct answer is: API

Question **255**

Correct

Mark 1.00 out of 1.00

You need to create an Azure resource to store data in Azure Table storage.
Which command should you run?

- a. az storage container create
- b. az storage share create
- c. az storage account create
- d. az cosmosdb create

Your answer is correct.

The correct answer is:
az storage account create

Question **256**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

Transcribing audio files is an example of [cognitive] analytics.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

Transcribing audio files is an example of [cognitive] analytics.

Question **257**

Correct

Mark 1.00 out of 1.00

You need to ensure that users use multi-factor authentication (MFA) when connecting to an Azure SQL database.
Which type of authentication should you use?

- a. SQL authentication
- b. service principal authentication
- c. Azure Active Directory (Azure AD) authentication
- d. certificate authentication

Your answer is correct.

The correct answer is:
Azure Active Directory (Azure AD) authentication

Question **258**

Correct

Mark 1.00 out of 1.00

What is a benefit of hosting a database on Azure SQL managed instance as compared to an Azure SQL database?

- a. support for encryption at rest
- b. native support for cross-database queries and transactions
- c. system-initiated automatic backups
- d. built-in high availability

Your answer is correct.

The correct answer is:
native support for cross-database queries and transactions

Question **259**

Incorrect

Mark 0.00 out of 1.00

What is the primary purpose of a data warehouse?

a. to provide read-only storage of relational and non-relational historical data

b. to provide storage for transactional line-of-business (LOB) applications

c. to provide answers to complex queries that rely on data from multiple sources

d. to provide transformation services between source and target data stores

Your answer is incorrect.

The correct answer is: **to provide answers to complex queries that rely on data from multiple sources**

**Explanation:**

A **data warehouse** is primarily designed to:

- Consolidate data from multiple sources
- Support **complex analytical queries** and reporting
- Enable business intelligence and decision-making

The correct answer is:
to provide answers to complex queries that rely on data from multiple sources

Question **260**

Correct

Mark 1.00 out of 1.00

Which storage solution supports access control lists (ACLs) at the file and folder level?

a. Azure Cosmos DB

b. Azure Data Lake Storage

c. Azure Queue storage

d. Azure Blob storage

Your answer is correct.

The correct answer is: Azure Data Lake Storage

Question **261**

Correct

Mark 1.00 out of 1.00

You have a database that stores data in the format shown in the following table.

Of which type of database is this an example?

| Key | Customer | | Product | |
|---|---|---|---|---|
| | Name | Address | Name | Price |
| 1000 | Ben Smith | 1 Main St. | Spark plug | 20.99 |
| 1001 | Carlos Grilo | 123 Elm Pl. | Spanner | 10.25 |

Of which type of database is this an example?

- a. key-value
- b. column family
- c. relational
- d. document

Your answer is correct.

The correct answer is: column family

Question **262**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Azure Blob Storage page blobs are optimized for random read and write operations.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

Azure Blob Storage page blobs are optimized for [random read and write operations.]

Question **263**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

A relational database must be used when strong consistency guarantees are required.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

A relational database must be used when [strong consistency guarantees are required.]

Question **264**

Correct

Mark 1.00 out of 1.00

What should you use to automatically delete blobs from Azure Blob Storage?

a. a lifecycle management policy

b. soft delete

c. archive storage

d. the change feed

Your answer is correct.

The correct answer is: a lifecycle management policy

Question **265**

Partially correct

Mark 1.00 out of 2.00

At which two levels can you set the throughput for an Azure Cosmos DB account? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [x] a. container
- [x] b. partition
- [ ] c. database
- [ ] d. item

Your answer is partially correct.

**Correct Answers: database, container**

**Explanation:**

In **Azure Cosmos DB**, you can configure **throughput (RU/s)** at the following two levels:

1. **Database Level (Shared Throughput):**
   - The specified throughput is **shared** among all containers in the database.
   - This is useful when you have multiple containers with unpredictable workloads.
2. **Container Level (Dedicated Throughput):**
   - Throughput is **dedicated** to a single container.
   - Ensures consistent performance for that specific container.

You have correctly selected 1.

The correct answers are:
database, container

Question **266**

Correct

Mark 1.00 out of 1.00

You need to gather real-time telemetry data from a mobile application.
Which type of workload describes this scenario?

- a. Online Transaction Processing (OLTP)
- b. massively parallel processing (MPP)
- c. streaming
- d. batch

Your answer is correct.

The correct answer is: streaming

Question **267**

Partially correct

Mark 2.00 out of 3.00

What are three characteristics of an Online Transaction Processing (OLTP) workload? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [ ] a. denormalized data
- [x] b. schema defined in a database
- [x] c. normalized data
- [ ] d. schema defined when reading unstructured data from a database
- [ ] e. light writes and heavy reads
- [x] f. heavy writes and moderate reads

Your answer is partially correct.

You have correctly selected 2.

The correct answers are:
denormalized data,

heavy writes and moderate reads,

schema defined in a database

Question **268**

Incorrect

Mark 0.00 out of 1.00

Select the answer that correctly completes the sentence.

In a relational database, each row in a table has a different set of columns

Your answer is incorrect.

**Correct Answer:** the same set of columns

**Reason:**

In a **relational database**, tables are structured with a **fixed schema**. This means:

- **All rows** (records) in a table follow the **same structure**,
- They have the **same set of columns**, though the **values** may vary.

The correct answer is:
Select the answer that correctly completes the sentence.
In a relational database, each row in a table has [ the same set of columns ]

Question **269**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

Microsoft Power BI can be used to simplify the process of embedding the results of data analysis on webpages for external users.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

[Microsoft Power BI] can be used to simplify the process of embedding the results of data analysis on webpages for external users.

Question **270**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

An online analytical processing (OLAP) model stores aggregated data that is optimized for analytical workloads.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

An online analytical processing (OLAP) model stores [aggregated] data that is optimized for analytical workloads.

Question **271**

Correct

Mark 3.00 out of 3.00

DRAG DROP -
Match the Azure Data Factory components to the appropriate descriptions.
To answer, drag the appropriate component from the column on the left to its description on the right. Each component may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.
Select and Place:

| | |
|---|---|
| Dataset | A representation of data structures within data stores |
| Linked service | The information used to connect to external resources |
| Pipeline | A logical grouping of activities that performs a unit of work and can be scheduled |

Mapping data flow

Your answer is correct.

The correct answer is:
DRAG DROP -
Match the Azure Data Factory components to the appropriate descriptions.
To answer, drag the appropriate component from the column on the left to its description on the right. Each component may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.
Select and Place:

[Dataset]A representation of data structures within data stores

[Linked service]The information used to connect to external resources

[Pipeline]A logical grouping of activities that performs a unit of work and can be scheduled

Question **272**

Incorrect

Mark 0.00 out of 1.00

Select the answer that correctly completes the sentence.

An Azure SQL Database elastic pool can be used as part of a lift-and-shift solution for migrating an on-premises application to the cloud with minimal changes.

Your answer is incorrect.

The correct answer is:
Select the answer that correctly completes the sentence.

[An instance of SQL Server on Azure Virtual Machines] can be used as part of a lift-and-shift solution for migrating an on-premises application to the cloud with minimal changes.

Question **273**

Correct

Mark 1.00 out of 1.00

Your company is designing a data store for internet-connected temperature sensors.
The collected data will be used to analyze temperature trends.
Which type of data store should you use?

- a. time series
- b. relational
- c. columnar
- d. graph

Your answer is correct.

The correct answer is:
time series

Question **274**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

A view is a virtual table that contains content defined by a query.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

[A view] is a virtual table that contains content defined by a query.

Question **275**

Partially correct

Mark 1.50 out of 3.00

Match the cloud service model to the appropriate Azure relational database service.

To answer, drag the appropriate cloud service model from the column on the left to its service on the right. Each model may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

| | | |
|---|---|---|
| Infrastructure as a service | ✗ | Azure SQL Database |
| Platform as a service | ✓ | Azure SQL Managed Instance |

[[]1] SQL Server on Azure virtual machine

Infrastructure as a service | Platform as a service | Software as a service

Your answer is partially correct.

You have correctly selected 1.

The correct answer is:
Match the cloud service model to the appropriate Azure relational database service.

To answer, drag the appropriate cloud service model from the column on the left to its service on the right. Each model may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

[Platform as a service] Azure SQL Database

[Platform as a service] Azure SQL Managed Instance

[[]1] SQL Server on Azure virtual machine

Question **276**

Correct

Mark 1.00 out of 1.00

You have a banking application that transfers money in to and out of accounts.

Of which type of solution is this is an example?

- a. online transaction processing (OLTP) ✓
- b. online analytical processing (OLAP)
- c. a data warehouse
- d. extract, transform and load (ETL)

Your answer is correct.

The correct answer is: online transaction processing (OLTP)

Question **277**

Correct

Mark 1.00 out of 1.00

Which Azure Cosmos DB API should you use for a graph database?

- a. Cassandra
- b. Table
- c. Gremlin
- d. Core (SQL)

Your answer is correct.

The correct answer is: Gremlin

Question **278**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

Relational data is stored in a tabular form of rows and columns.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

Relational data is stored in [a tabular form of rows and columns.]

Question **279**

Correct

Mark 3.00 out of 3.00

DRAG DROP -
Match the Azure Cosmos DB APIs to the appropriate data structures.
To answer, drag the appropriate API from the column on the left to its data structure on the right. Each API may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

| | |
|---|---|
| Gremlin API | Graph data |
| MongoDB API | JSON documents |
| Table API | Key/value data |
| Cassandra API | |

Your answer is correct.

The correct answer is:
DRAG DROP -
Match the Azure Cosmos DB APIs to the appropriate data structures.
To answer, drag the appropriate API from the column on the left to its data structure on the right. Each API may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

[Gremlin API]Graph data

[MongoDB API ]JSON documents

[Table API]Key/value data

Question **280**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

To create a dashboard, you can use either the Microsoft Power BI Desktop app or the Power BI service

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

To create a dashboard, you can use [either the Microsoft Power BI Desktop app or the Power BI service]

Question **281**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

A database administrator is responsible for managing the security of the data in a database.

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

[A database administrator] is responsible for managing the security of the data in a database.

Question **282**

Correct

Mark 2.00 out of 2.00

Which two features distinguish Delta Lake from Azure Data Lake Storage? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- [ ] a. support for batch data
- [x] b. transactional consistency
- [ ] c. support for an Apache Spark runtime
- [ ] d. support for streaming data
- [x] e. schema enforcement

Your answer is correct.

The correct answers are: schema enforcement, transactional consistency

Question **283**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

In a transactional database that guarantees atomicity, all statements in a transaction must complete successfully or the transaction must roll back. Yes

In a transactional database that guarantees consistency, all data changes will be logged and accounted for. No

In a transactional database that guarantees isolation, processes can see only data in a consistent state. Yes

Your answer is correct.

The correct answer is:
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

In a transactional database that guarantees atomicity, all statements in a transaction must complete successfully or the transaction must roll back. [Yes]
In a transactional database that guarantees consistency, all data changes will be logged and accounted for. [No]
In a transactional database that guarantees isolation, processes can see only data in a consistent state. [Yes]

Question **284**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completed the sentence.

Scatter charts are used to compare two numeric measures and identify a relationship or correlation between them.

Your answer is correct.

The correct answer is:
Select the answer that correctly completed the sentence.

[Scatter charts] are used to compare two numeric measures and identify a relationship or correlation between them.

Question **285**

Correct

Mark 3.00 out of 3.00

DRAG DROP -

Match the types of visualizations to the appropriate descriptions.

To answer, drag the appropriate visualization type from the column on the left to its description on the right. Each visualization type may be used once, more than once, or not at all.

Treemap A chart of colored, nested rectangles that displays individual data points represented by the size and color of a relative rectangle.

Key influencer A chart that displays the major contributors of a selected result or value.

Scatter A chart that shows the relationship between two numerical values.

Select and Place:

Your answer is correct.

The correct answer is:

DRAG DROP -

Match the types of visualizations to the appropriate descriptions.

To answer, drag the appropriate visualization type from the column on the left to its description on the right. Each visualization type may be used once, more than once, or not at all.

[Treemap]A chart of colored, nested rectangles that displays individual data points represented by the size and color of a relative rectangle.

[Key influencer]A chart that displays the major contributors of a selected result or value.

[Scatter]A chart that shows the relationship between two numerical values.

Select and Place:

Question **286**

Correct

Mark 1.00 out of 1.00

What is required to provision Azure Data Lake Storage in an Azure Storage account?

a. Hierarchical namespace must be disabled.

b. Hierarchical namespace must be enabled.

c. Versioning must be disabled.

d. Versioning must be enabled.

Your answer is correct.

The correct answer is: Hierarchical namespace must be enabled.

Question **287**

Correct

Mark 1.00 out of 1.00

You have a dedicated SQL pool in Azure Synapse Analytics that is only used actively every night for eight hours.
You need to minimize the cost of the dedicated SQL pool as much as possible during idle times. The solution must ensure that the data remains intact.
What should you do on the dedicated SQL pool?

- a. Delete the pool
- b. Create a user-defined restore point.
- c. Scale down the data warehouse units (DWUs).
- d. Pause the pool.

Your answer is correct.

The correct answer is: Pause the pool.

Question **288**

Incorrect

Mark 0.00 out of 3.00

DRAG DROP -
Match the Azure services to the appropriate requirements.
To answer, drag the appropriate service from the column on the left to its requirement on the right. Each service may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Azure Synapse Analytics ⊗ Output data to Parque format
[ [1]]Store data that is in Parquet format

Azure SQL Database ⊗ Persist a tabular representation of data that is stored in Parquet format

Azure Data Lake Storage

Your answer is incorrect.

**Explanation:**

**Output data to Parquet format → Azure Synapse Analytics**

Synapse supports **exporting query results** to Parquet format via **PolyBase or CETAS (CREATE EXTERNAL TABLE AS SELECT)**.

**Store data that is in Parquet format → Azure Data Lake Storage**

Data Lake is ideal for storing **structured, semi-structured, or unstructured files**, including **Parquet**, **CSV**, **JSON**, etc.

**Persist a tabular representation of data stored in Parquet format → Azure Synapse Analytics**

Synapse can **query external Parquet files using external tables** and create **materialized views or internal tables** to represent that data.

The correct answer is:
DRAG DROP -
Match the Azure services to the appropriate requirements.
To answer, drag the appropriate service from the column on the left to its requirement on the right. Each service may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

[Azure SQL Database]Output data to Parque format
[ [1]]Store data that is in Parquet format

[Azure Synapse Analytics] Persist a tabular representation of data that is stored in Parquet format

Question **289**

Incorrect

Mark 0.00 out of 1.00

After you create Microsoft Power BI data models and reports, you can publish them to Power BI Desktop ✗ .

Your answer is incorrect.

The correct answer is:
After you create Microsoft Power BI data models and reports, you can publish them to [the Power BI service].

Question **290**

Correct

Mark 3.00 out of 3.00

Match the job roles to the appropriate tasks.

To answer, drag the appropriate role from the column on the left to its task on the right. Each role may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

| | | |
|---|---|---|
| Database administrator | ✓ | Database backup and recovery |
| Data analyst | ✓ | Enabling advanced analytic capabilities through reports |
| Data Engineer | ✓ | Monitoring data pipelines |

Your answer is correct.

The correct answer is:
Match the job roles to the appropriate tasks.

To answer, drag the appropriate role from the column on the left to its task on the right. Each role may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

[Database administrator] Database backup and recovery
[Data analyst] Enabling advanced analytic capabilities through reports
[Data Engineer] Monitoring data pipelines

Question **291**

Correct

Mark 1.00 out of 1.00

Which Azure Blob Storage feature enables you to sync a Blob Storage account to an external cache?

- a. versioning
- b. change feed
- c. the Hot access tier
- d. snapshots

Your answer is correct.

The correct answer is: change feed

Question **292**

Correct

Mark 2.00 out of 2.00

You have an Azure Cosmos DB account that uses the Core (SQL) API.
Which two settings can you configure at the container level? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [ ] a. the API
- [x] b. the partition key
- [ ] c. the read region
- [x] d. the throughput

Your answer is correct.

The correct answers are: the throughput, the partition key

Question **293**

Correct

Mark 3.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| You can use Azure Data Studio to query a Microsoft SQL Server big data cluster. | YES |
| You can use Microsoft SQL Server Management Studio (SSMS) to query an Azure Synapse Analytics data warehouse. | YES |
| You can use MySQL Workbench to query Azure Database for MariaDB databases. | YES |

Your answer is correct.

The correct answer is: You can use Azure Data Studio to query a Microsoft SQL Server big data cluster. → YES, You can use Microsoft SQL Server Management Studio (SSMS) to query an Azure Synapse Analytics data warehouse. → YES, You can use MySQL Workbench to query Azure Database for MariaDB databases. → YES

Question **294**

Correct

Mark 3.00 out of 3.00

DRAG DROP -
Match the types of workloads to the appropriate scenarios.
To answer, drag the appropriate workload type from the column on the left to its scenario on the right. Each workload type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

| | |
|---|---|
| Batch | Data for a product catalog will be loaded every 12 hours to a data warehouse. |
| Streaming | Thousands of data sets per second for online purchases will be loaded into a data warehouse in real time |
| Batch | Updates to inventory data will be loaded to a data warehouse every 1 million transactions. |

Select and Place:

Batch

Your answer is correct.

The correct answer is:
DRAG DROP -
Match the types of workloads to the appropriate scenarios.
To answer, drag the appropriate workload type from the column on the left to its scenario on the right. Each workload type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

[Batch ]Data for a product catalog will be loaded every 12 hours to a data warehouse.

[Streaming]Thousands of data sets per second for online purchases will be loaded into a data warehouse in real time

[Batch ]Updates to inventory data will be loaded to a data warehouse every 1 million transactions.

Select and Place:

Question **295**

Correct

Mark 2.00 out of 2.00

What are two uses of data visualization? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

- [ ] a. Implement machine learning to predict future values
- [x] b. Represent trends and patterns over time
- [ ] c. Enforce business logic across reports
- [x] d. Communicate the significance of data

Your answer is correct.

The correct answers are:
Represent trends and patterns over time,

Communicate the significance of data

Question **296**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completes the sentence.

In a graph database, edges specify relationships between vertexes .

Your answer is correct.

The correct answer is:
Select the answer that correctly completes the sentence.

In a graph database, edges specify relationships between [vertexes].

Question **297**

Correct

Mark 3.00 out of 3.00

Match the processes to the appropriate scenarios.
To answer, drag the appropriate process from the column on the left to its scenario on the right. Each process may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Data is retrieved and processed before being saved to a target system | Extract, transform, and load (ETL) processing

Data is saved to a target system, and then processed | Extract, load, and transform (ELT) processing

Data is extracted from a source system, and then saved to a target system | Data ingestion

Your answer is correct.

The correct answer is:
Match the processes to the appropriate scenarios.
To answer, drag the appropriate process from the column on the left to its scenario on the right. Each process may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.

Data is retrieved and processed before being saved to a target system [Extract, transform, and load (ETL) processing]
Data is saved to a target system, and then processed [Extract, load, and transform (ELT) processing]
Data is extracted from a source system, and then saved to a target system [Data ingestion]

Question **298**

Correct

Mark 1.00 out of 1.00

Structured data where each row represents a single data entity uses which type of schema?

a. tabular
b. blob
c. XML
d. JSON

Your answer is correct.

The correct answer is: tabular

Question **299**

Correct

Mark 1.00 out of 1.00

Select the answer that correctly completed the sentence.

The pipeline engine used in Azure Data Factory is also available in Azure Synapse Analytics. .

Your answer is correct.

The correct answer is:
Select the answer that correctly completed the sentence.

The pipeline engine used in Azure Data Factory is also available in [Azure Synapse Analytics.].

Question **300**

Correct

Mark 1.00 out of 1.00

To complete the sentence, select the appropriate option in the answer area.

To configure an Azure Storage account to support both security at the folder level and atomic directory manipulation,

enable the hierarchical namespace.

Your answer is correct.

The correct answer is:
To complete the sentence, select the appropriate option in the answer area.

To configure an Azure Storage account to support both security at the folder level and atomic directory manipulation,[enable the hierarchical namespace. ]

Question **301**

Correct

Mark 1.00 out of 1.00

How is raw data stored in a data lakehouse?

- a. as files
- b. as queues
- c. as tables
- d. as disks

Your answer is correct.

The correct answer is: as files

Question **302**

Correct

Mark 1.00 out of 1.00

A data analyst is responsible for finding hidden patterns in data and creating graphical reports that deliver information in a useful and appealing format.

Your answer is correct.

The correct answer is:
[A data analyst] is responsible for finding hidden patterns in data and creating graphical reports that deliver information in a useful and appealing format.

Question **303**

Incorrect

Mark 0.00 out of 1.00

Which Azure service provides the highest compatibility for databases migrated from Microsoft SQL Server 2019 Enterprise edition?

a. Azure Database for MySQL

b. Azure SQL Managed Instance

c. an Azure SQL Database elastic pool

d. Azure SQL Database

Your answer is incorrect.

**Correct Answer**: Azure SQL Managed Instance

**Explanation:**

**Azure SQL Managed Instance** provides the **highest compatibility** with **Microsoft SQL Server 2019 Enterprise Edition**, including:

- **Cross-database queries**
- **SQL Agent**
- **CLR, Service Broker, and more enterprise features**
- Near **100% compatibility** with the on-premises SQL Server engine

This makes it the best choice for migrating existing SQL Server workloads with minimal changes.

The correct answer is:
Azure SQL Managed Instance

Question **304**

Partially correct

Mark 1.00 out of 3.00

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.

| | |
|---|---|
| The Azure Cosmos DB API is configured separately for each database in an Azure Cosmos DB account. | YES |
| Partition keys are used in Azure Cosmos DB to optimize queries. | YES |
| Items contained in the same Azure Cosmos DB logical partition can have different partition keys. | YES |

Your answer is partially correct.

| Statement | Answer | Explanation |
|---|---|---|
| **The Azure Cosmos DB API is configured separately for each database in an Azure Cosmos DB account.** | **No** | The API (e.g., Core/SQL, MongoDB, Cassandra, Gremlin, Table) is selected **at the account level**, not per database. All databases and containers within that Cosmos DB account use the same API. |
| **Partition keys are used in Azure Cosmos DB to optimize queries.** | **Yes** | Partition keys help distribute data across physical partitions and improve **query performance and scalability**. |
| **Items contained in the same Azure Cosmos DB logical partition can have different partition keys.** | **No** | All items in the **same logical partition** must share the **same partition key value**. That's how partitions are defined. |

You have correctly selected 1.

The correct answer is: The Azure Cosmos DB API is configured separately for each database in an Azure Cosmos DB account. → NO, Partition keys are used in Azure Cosmos DB to optimize queries. → YES, Items contained in the same Azure Cosmos DB logical partition can have different partition keys. → NO